

## Breakout-3

Amna

Samreen

Ghazia

Misha

Hamid

Haris

Abdullah

```
In [1]: from google.colab import drive  
drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force\_remount=True).


```
In [2]: import numpy as np  
import matplotlib.pyplot as plt  
import seaborn as sns  
import pandas as pd
```

```
In [3]: df = pd.read_csv('/content/drive/MyDrive/weatherAUS.csv')  
df
```

Out[3]:

	Date	Location	MinTemp	MaxTemp	Rainfall	Evaporation	Sunshine	WindGustI
0	2008-12-01	Albury	13.4	22.9	0.6	NaN	NaN	
1	2008-12-02	Albury	7.4	25.1	0.0	NaN	NaN	WN
2	2008-12-03	Albury	12.9	25.7	0.0	NaN	NaN	WS
3	2008-12-04	Albury	9.2	28.0	0.0	NaN	NaN	
4	2008-12-05	Albury	17.5	32.3	1.0	NaN	NaN	
...	...	...	...	...	...	...	...	
145455	2017-06-21	Uluru	2.8	23.4	0.0	NaN	NaN	
145456	2017-06-22	Uluru	3.6	25.3	0.0	NaN	NaN	NN
145457	2017-06-23	Uluru	5.4	26.9	0.0	NaN	NaN	
145458	2017-06-24	Uluru	7.8	27.0	0.0	NaN	NaN	
145459	2017-06-25	Uluru	14.9	NaN	0.0	NaN	NaN	Ni

145460 rows × 23 columns



In [4]:

```
df.head()
```

Out[4]:

	Date	Location	MinTemp	MaxTemp	Rainfall	Evaporation	Sunshine	WindGustDir	W
0	2008-12-01	Albury	13.4	22.9	0.6	NaN	NaN		W
1	2008-12-02	Albury	7.4	25.1	0.0	NaN	NaN		WNW
2	2008-12-03	Albury	12.9	25.7	0.0	NaN	NaN		WSW
3	2008-12-04	Albury	9.2	28.0	0.0	NaN	NaN		NE
4	2008-12-05	Albury	17.5	32.3	1.0	NaN	NaN		W

5 rows × 23 columns

In [5]:

df.shape

Out[5]:

(145460, 23)

In [6]:

df.describe()

Out[6]:


	MinTemp	MaxTemp	Rainfall	Evaporation	Sunshine	WindGust
count	143975.000000	144199.000000	142199.000000	82670.000000	75625.000000	135197.0
mean	12.194034	23.221348	2.360918	5.468232	7.611178	40.0
std	6.398495	7.119049	8.478060	4.193704	3.785483	13.0
min	-8.500000	-4.800000	0.000000	0.000000	0.000000	6.0
25%	7.600000	17.900000	0.000000	2.600000	4.800000	31.0
50%	12.000000	22.600000	0.000000	4.800000	8.400000	39.0
75%	16.900000	28.200000	0.800000	7.400000	10.600000	48.0
max	33.900000	48.100000	371.000000	145.000000	14.500000	135.0

In [7]:

df.describe(include = ['object'])

Out[7]:

	Date	Location	WindGustDir	WindDir9am	WindDir3pm	RainToday	RainTomor
<b>count</b>	145460	145460	135134	134894	141232	142199	142
<b>unique</b>	3436	49	16	16	16	2	
<b>top</b>	2013-11-12	Canberra	W	N	SE	No	
<b>freq</b>	49	3436	9915	11758	10838	110319	110


In [8]: `df.isnull().sum()`

```
Out[8]: Date          0
Location          0
MinTemp         1485
MaxTemp         1261
Rainfall         3261
Evaporation     62790
Sunshine        69835
WindGustDir     10326
WindGustSpeed   10263
WindDir9am      10566
WindDir3pm       4228
WindSpeed9am    1767
WindSpeed3pm    3062
Humidity9am     2654
Humidity3pm     4507
Pressure9am     15065
Pressure3pm     15028
Cloud9am        55888
Cloud3pm        59358
Temp9am         1767
Temp3pm         3609
RainToday       3261
RainTomorrow    3267
dtype: int64
```

In [9]: `df.duplicated().sum()`

Out[9]: 0

```
In [10]: missing_percentages = df.isnull().mean()*100
missing_percentages
```

```
Out[10]: Date          0.000000
Location          0.000000
MinTemp           1.020899
MaxTemp           0.866905
Rainfall          2.241853
Evaporation       43.166506
Sunshine         48.009762
WindGustDir       7.098859
WindGustSpeed     7.055548
WindDir9am        7.263853
WindDir3pm        2.906641
WindSpeed9am      1.214767
WindSpeed3pm      2.105046
Humidity9am       1.824557
Humidity3pm       3.098446
Pressure9am       10.356799
Pressure3pm       10.331363
Cloud9am          38.421559
Cloud3pm          40.807095
Temp9am           1.214767
Temp3pm           2.481094
RainToday         2.241853
RainTomorrow      2.245978
dtype: float64
```

```
In [11]: df.columns
```

```
Out[11]: Index(['Date', 'Location', 'MinTemp', 'MaxTemp', 'Rainfall', 'Evaporation',
               'Sunshine', 'WindGustDir', 'WindGustSpeed', 'WindDir9am', 'WindDir3pm',
               'WindSpeed9am', 'WindSpeed3pm', 'Humidity9am', 'Humidity3pm',
               'Pressure9am', 'Pressure3pm', 'Cloud9am', 'Cloud3pm', 'Temp9am',
               'Temp3pm', 'RainToday', 'RainTomorrow'],
              dtype='object')
```

Imputing Missing values in RainTomorrow column using mode.

```
In [12]: df['RainTomorrow'].value_counts()
```

```
Out[12]: RainTomorrow
No      110316
Yes      31877
Name: count, dtype: int64
```

```
In [13]: df['RainTomorrow'].unique()
```

```
Out[13]: array(['No', 'Yes', nan], dtype=object)
```

```
In [14]: mode_value = df['RainTomorrow'].mode()[0]
print(mode_value)
```

No

```
In [15]: df['RainTomorrow'].fillna(value=mode_value, inplace=True)
```

```
In [16]: df[['RainTomorrow']]
```

Out[16]:

RainTomorrow	
0	No
1	No
2	No
3	No
4	No
...	...
145455	No
145456	No
145457	No
145458	No
145459	No

145460 rows × 1 columns

In [17]: `df.isnull().sum()`

```
Out[17]: Date          0
Location          0
MinTemp         1485
MaxTemp         1261
Rainfall         3261
Evaporation     62790
Sunshine        69835
WindGustDir     10326
WindGustSpeed   10263
WindDir9am      10566
WindDir3pm       4228
WindSpeed9am    1767
WindSpeed3pm    3062
Humidity9am     2654
Humidity3pm     4507
Pressure9am     15065
Pressure3pm     15028
Cloud9am        55888
Cloud3pm        59358
Temp9am         1767
Temp3pm         3609
RainToday       3261
RainTomorrow     0
dtype: int64
```

Imputing missing values in Rain Fall column using Median

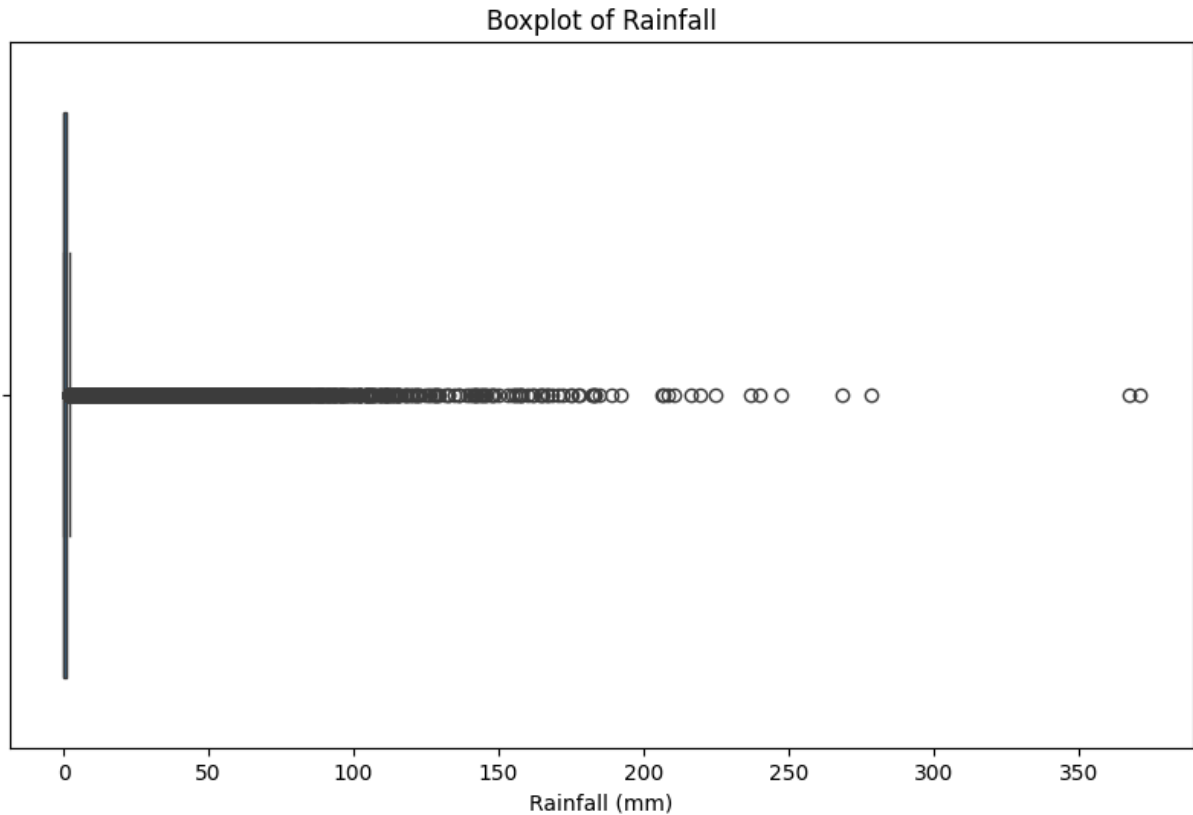
In [18]: `df['Rainfall'].describe()`

```
Out[18]: count    142199.000000
         mean       2.360918
         std       8.478060
         min        0.000000
         25%        0.000000
         50%        0.000000
         75%        0.800000
         max       371.000000
         Name: Rainfall, dtype: float64
```

```
In [19]: df['Rainfall'].isnull().sum()
```

```
Out[19]: 3261
```

```
In [20]: plt.figure(figsize=(10, 6))
         sns.boxplot(x=df['Rainfall'])
         plt.title('Boxplot of Rainfall')
         plt.xlabel('Rainfall (mm)')
         plt.show()
```



```
In [21]: Q1 = df['Rainfall'].quantile(0.25)
         Q3 = df['Rainfall'].quantile(0.75)
         IQR = Q3 - Q1
         lower_bound = Q1 - 1.5 * IQR
         upper_bound = Q3 + 1.5 * IQR

         outliers = df[(df['Rainfall'] < lower_bound) | (df['Rainfall'] > upper_bound)]
         percentage_outliers = len(outliers) / len(df) * 100
         print(f"Percentage of outliers in 'Rainfall': {percentage_outliers:.2f}%")
```

Percentage of outliers in 'Rainfall': 17.58%

```
In [22]: median_value = df['Rainfall'].median()
```

```
In [23]: print(median_value)
```

0.0

```
In [24]: df['Rainfall'].fillna(value = median_value, inplace = True)
```

```
In [25]: df['Rainfall'].isnull().sum()
```

Out[25]: 0

```
In [26]: df.isnull().sum()
```

```
Out[26]: Date                0
Location                  0
MinTemp                 1485
MaxTemp                 1261
Rainfall                 0
Evaporation             62790
Sunshine                69835
WindGustDir             10326
WindGustSpeed           10263
WindDir9am              10566
WindDir3pm              4228
WindSpeed9am            1767
WindSpeed3pm            3062
Humidity9am             2654
Humidity3pm             4507
Pressure9am             15065
Pressure3pm             15028
Cloud9am                55888
Cloud3pm                59358
Temp9am                 1767
Temp3pm                 3609
RainToday                3261
RainTomorrow            0
dtype: int64
```

Imputing missing values in Humidity9am column using Mean

```
In [27]: df[['Humidity9am']]
```



Out[27]:

Humidity9am	
0	71.0
1	44.0
2	38.0
3	45.0
4	82.0
...	...
145455	51.0
145456	56.0
145457	53.0
145458	51.0
145459	62.0

145460 rows × 1 columns

In [28]: `df['Humidity9am'].describe()`

Out[28]:

count	142806.000000
mean	68.880831
std	19.029164
min	0.000000
25%	57.000000
50%	70.000000
75%	83.000000
max	100.000000

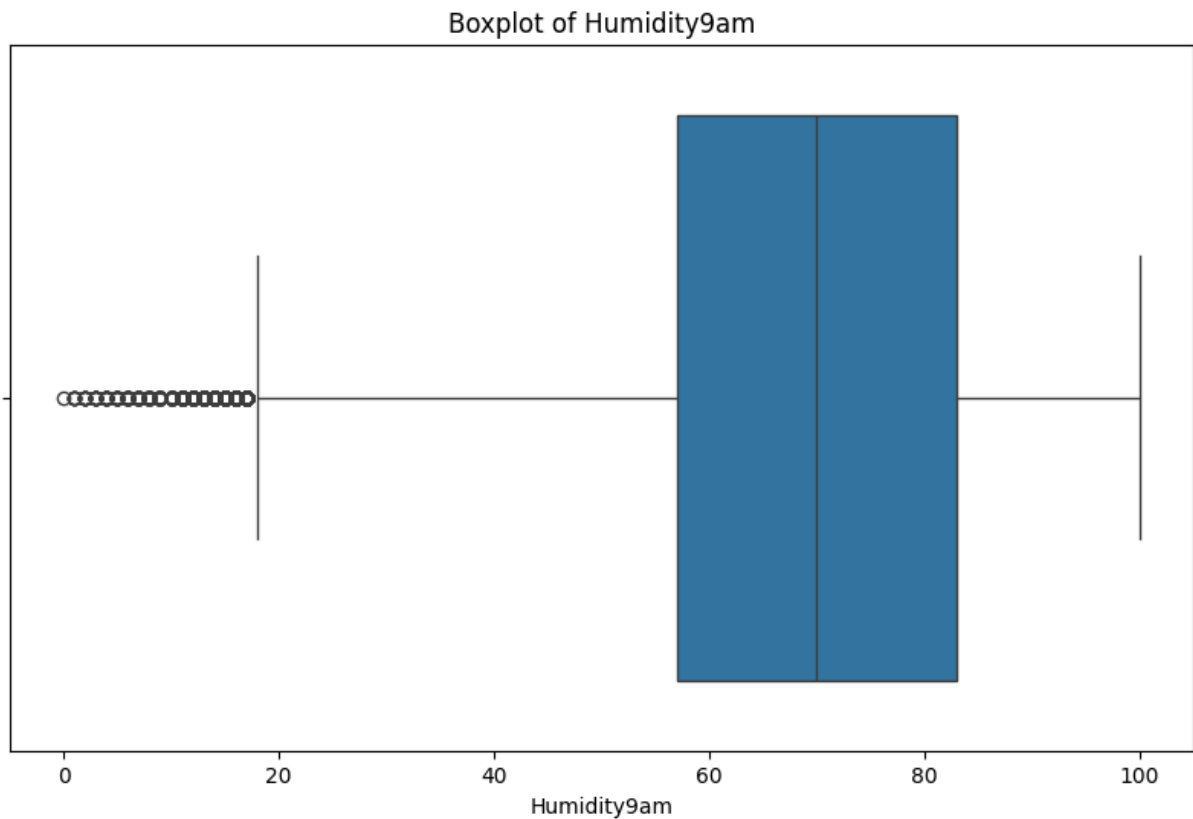
Name: Humidity9am, dtype: float64

In [29]: `df['Humidity9am'].isnull().sum()`

Out[29]: 2654

In [30]:

```
plt.figure(figsize=(10, 6))
sns.boxplot(x=df['Humidity9am'])
plt.title('Boxplot of Humidity9am')
plt.xlabel('Humidity9am')
plt.show()
```



```
In [31]: Q1 = df['Humidity9am'].quantile(0.25)
Q3 = df['Humidity9am'].quantile(0.75)
IQR = Q3 - Q1
lower_bound = Q1 - 1.5 * IQR
upper_bound = Q3 + 1.5 * IQR

outliers = df[(df['Humidity9am'] < lower_bound) | (df['Humidity9am'] > upper_bound)]
percentage_outliers = len(outliers) / len(df) * 100
print(f"Percentage of outliers in 'Humidity9am': {percentage_outliers:.2f}%")
```

Percentage of outliers in 'Humidity9am': 0.98%

```
In [32]: mean_Humidity9am = df['Humidity9am'].mean()
```

```
In [33]: print(mean_Humidity9am)
```

68.88083133761887

```
In [34]: df['Humidity9am'].fillna( value = mean_Humidity9am, inplace = True)
```

```
In [35]: df.isnull().sum()
```

```
Out[35]: Date          0
         Location      0
         MinTemp      1485
         MaxTemp      1261
         Rainfall      0
         Evaporation  62790
         Sunshine     69835
         WindGustDir   10326
         WindGustSpeed 10263
         WindDir9am    10566
         WindDir3pm    4228
         WindSpeed9am  1767
         WindSpeed3pm  3062
         Humidity9am   0
         Humidity3pm   4507
         Pressure9am   15065
         Pressure3pm   15028
         Cloud9am      55888
         Cloud3pm      59358
         Temp9am       1767
         Temp3pm       3609
         RainToday     3261
         RainTomorrow  0
         dtype: int64
```

Imputing missing values WindGustSpeed column using median

```
In [36]: df[['WindGustSpeed']]
```

Out[36]:

	WindGustSpeed
0	44.0
1	44.0
2	46.0
3	24.0
4	41.0
...	...
145455	31.0
145456	22.0
145457	37.0
145458	28.0
145459	NaN

145460 rows × 1 columns

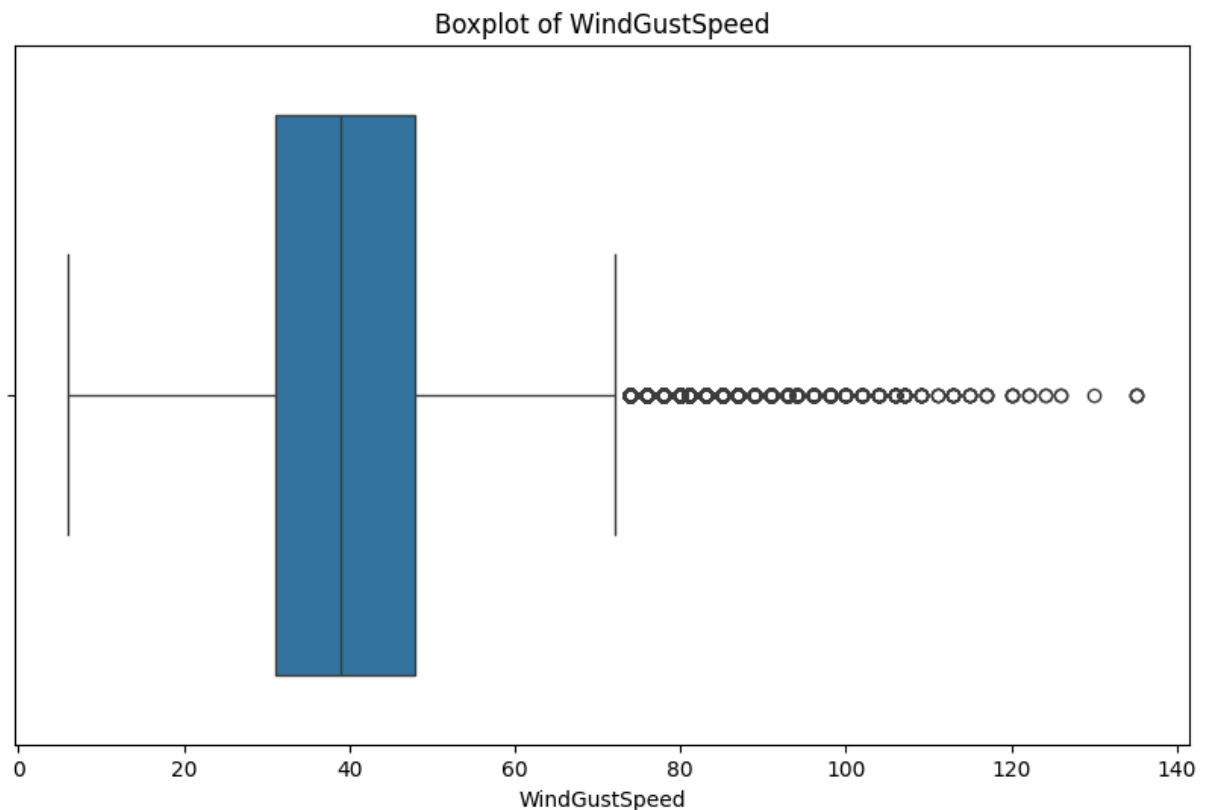
```
In [37]: df['WindGustSpeed'].describe()
```

```
Out[37]: count    135197.000000
mean         40.035230
std          13.607062
min           6.000000
25%          31.000000
50%          39.000000
75%          48.000000
max          135.000000
Name: WindGustSpeed, dtype: float64
```

```
In [38]: df['WindGustSpeed'].isnull().sum()
```

```
Out[38]: 10263
```

```
In [39]: plt.figure(figsize=(10, 6))
sns.boxplot(x=df['WindGustSpeed'])
plt.title('Boxplot of WindGustSpeed')
plt.xlabel('WindGustSpeed')
plt.show()
```



```
In [40]: Q1 = df['WindGustSpeed'].quantile(0.25)
Q3 = df['WindGustSpeed'].quantile(0.75)
IQR = Q3 - Q1
lower_bound = Q1 - 1.5 * IQR
upper_bound = Q3 + 1.5 * IQR

outliers = df[(df['WindGustSpeed'] < lower_bound) | (df['WindGustSpeed'] > upper_bound)]
```

```
percentage_outliers = len(outliers) / len(df) * 100  
print(f"Percentage of outliers in 'WindGustSpeed': {percentage_outliers:.2f}%")
```

Percentage of outliers in 'WindGustSpeed': 2.13%

```
In [41]: median_wgs = df['WindGustSpeed'].median()  
df['WindGustSpeed'].fillna(value = median_wgs, inplace = True)
```

```
In [42]: df.isnull().sum()
```

```
Out[42]: Date                0  
Location                  0  
MinTemp                 1485  
MaxTemp                 1261  
Rainfall                0  
Evaporation            62790  
Sunshine               69835  
WindGustDir            10326  
WindGustSpeed           0  
WindDir9am             10566  
WindDir3pm             4228  
WindSpeed9am           1767  
WindSpeed3pm           3062  
Humidity9am             0  
Humidity3pm            4507  
Pressure9am            15065  
Pressure3pm            15028  
Cloud9am               55888  
Cloud3pm               59358  
Temp9am                1767  
Temp3pm                3609  
RainToday              3261  
RainTomorrow           0  
dtype: int64
```

## Imputing missing values in Pressure9am column using median

```
In [43]: df[['Pressure9am']]
```

Out[43]:

Pressure9am	
0	1007.7
1	1010.6
2	1007.6
3	1017.6
4	1010.8
...	...
145455	1024.6
145456	1023.5
145457	1021.0
145458	1019.4
145459	1020.2

145460 rows × 1 columns

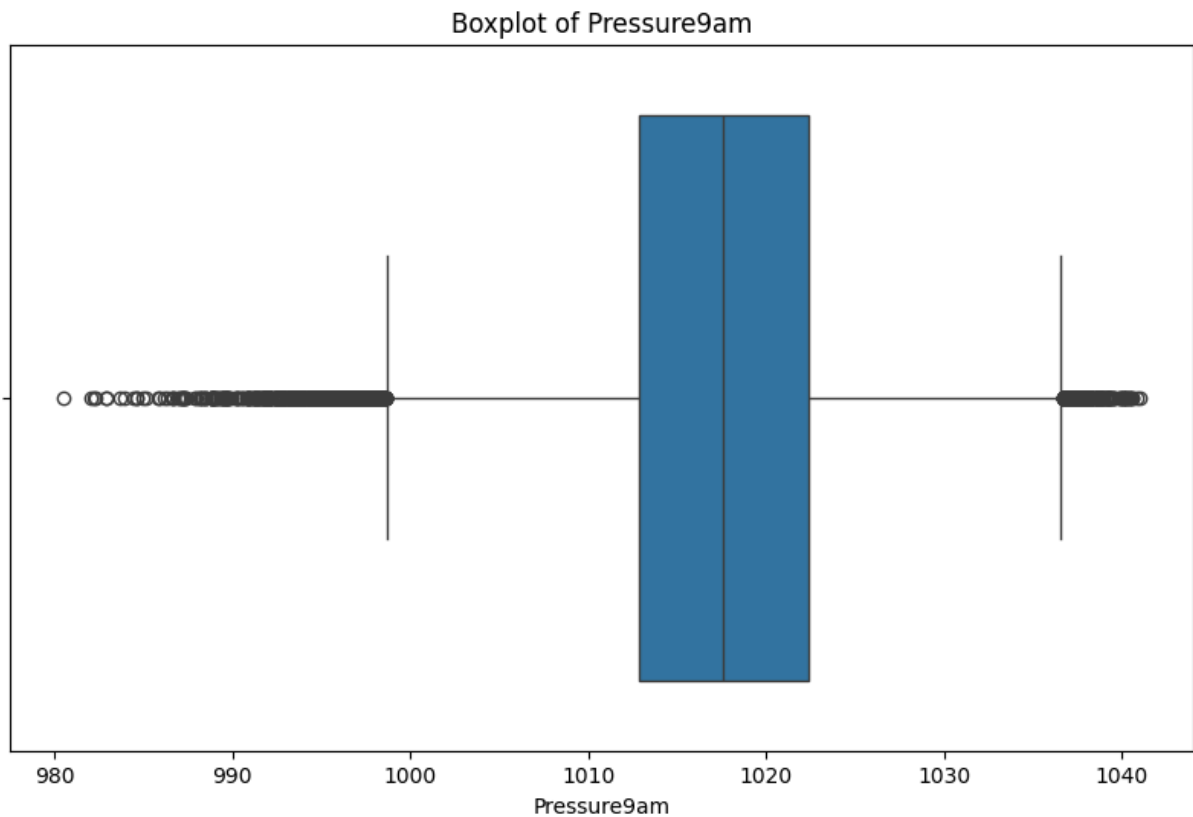
In [44]: `df['Pressure9am'].describe()`

```
Out[44]: count    130395.00000
mean      1017.64994
std        7.10653
min        980.50000
25%       1012.90000
50%       1017.60000
75%       1022.40000
max       1041.00000
Name: Pressure9am, dtype: float64
```

In [45]: `df['Pressure9am'].isnull().sum()`

Out[45]: 15065

```
In [46]: plt.figure(figsize=(10, 6))
sns.boxplot(x=df['Pressure9am'])
plt.title('Boxplot of Pressure9am')
plt.xlabel('Pressure9am')
plt.show()
```



```
In [47]: def outlier_per (df, column_name):
    """
    Calculate and return the percentage of outliers in a specified column of a Data
    Parameters:
        df (pd.DataFrame): The DataFrame containing the data.
        column_name (str): The name of the column to analyze for outliers.
    Returns:
        float: The percentage of values in the column that are considered outliers.
    """
    # Calculate the interquartile range (IQR)
    Q1 = df[column_name].quantile(0.25)
    Q3 = df[column_name].quantile(0.75)
    IQR = Q3 - Q1

    # Determine the bounds for outliers
    lower_bound = Q1 - 1.5 * IQR
    upper_bound = Q3 + 1.5 * IQR

    # Identify outliers
    outliers = df[(df[column_name] < lower_bound) | (df[column_name] > upper_bound)]

    # Calculate the percentage of outliers
    percentage_outliers = len(outliers) / len(df) * 100

    return percentage_outliers
outlier_per(df, 'Pressure9am')
```

Out[47]: 0.818781795682662

```
In [48]: Q1 = df['Pressure9am'].quantile(0.25)
Q3 = df['Pressure9am'].quantile(0.75)
IQR = Q3 - Q1
lower_bound = Q1 - 1.5 * IQR
upper_bound = Q3 + 1.5 * IQR

outliers = df[(df['Pressure9am'] < lower_bound) | (df['Pressure9am'] > upper_bound)]
percentage_outliers = len(outliers) / len(df) * 100
print(f"Percentage of outliers in 'Pressure9am': {percentage_outliers:.2f}%")
```

Percentage of outliers in 'Pressure9am': 0.82%

```
In [49]: median_pressure = df['Pressure9am'].median()
print(median_pressure)
```

1017.6

```
In [50]: df['Pressure9am'].fillna(value = median_pressure, inplace = True)
```

```
In [51]: df.isnull().sum()
```

```
Out[51]: Date                0
Location                  0
MinTemp                  1485
MaxTemp                  1261
Rainfall                  0
Evaporation             62790
Sunshine                 69835
WindGustDir              10326
WindGustSpeed             0
WindDir9am              10566
WindDir3pm               4228
WindSpeed9am             1767
WindSpeed3pm             3062
Humidity9am               0
Humidity3pm              4507
Pressure9am               0
Pressure3pm              15028
Cloud9am                  55888
Cloud3pm                  59358
Temp9am                   1767
Temp3pm                   3609
RainToday                 3261
RainTomorrow              0
dtype: int64
```

Covert RainToday and RainTomorrow columns into numerical column using label encoding method these were catergoical columns

```
In [52]: df.describe(include = ['object'])
```



Out[52]:

	Date	Location	WindGustDir	WindDir9am	WindDir3pm	RainToday	RainTomor
<b>count</b>	145460	145460	135134	134894	141232	142199	145
<b>unique</b>	3436	49	16	16	16	2	
<b>top</b>	2013-11-12	Canberra	W	N	SE	No	
<b>freq</b>	49	3436	9915	11758	10838	110319	113

In [53]: `from sklearn.preprocessing import LabelEncoder`In [54]: `le = LabelEncoder()`In [55]: `df['RainToday']=le.fit_transform(df['RainToday'])`In [56]: `df['RainTomorrow']=le.fit_transform(df['RainTomorrow'])`In [57]: `df[['RainToday','RainTomorrow']]`

Out[57]:

	RainToday	RainTomorrow
<b>0</b>	0	0
<b>1</b>	0	0
<b>2</b>	0	0
<b>3</b>	0	0
<b>4</b>	0	0
<b>...</b>	...	...
<b>145455</b>	0	0
<b>145456</b>	0	0
<b>145457</b>	0	0
<b>145458</b>	0	0
<b>145459</b>	0	0

145460 rows × 2 columns

Imputing Missing Values in Cloud9am USING mean

In [58]: `df[['Cloud9am']]`

Out[58]:

Cloud9am	
0	8.0
1	NaN
2	NaN
3	NaN
4	7.0
...	...
145455	NaN
145456	NaN
145457	NaN
145458	3.0
145459	8.0

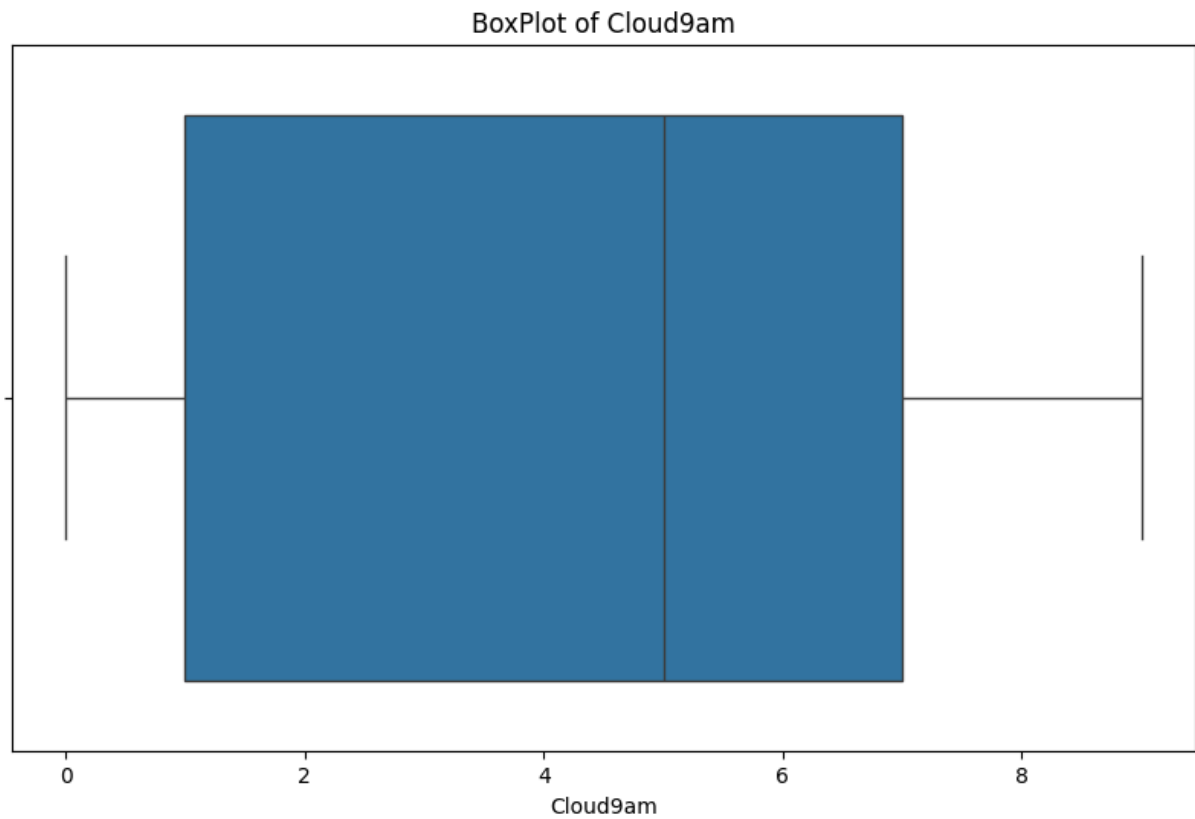
145460 rows × 1 columns

In [59]: `df['Cloud9am'].isnull().sum()`

Out[59]: 55888

In [60]: 

```
plt.figure(figsize=(10,6))
sns.boxplot(x=df['Cloud9am'])
plt.title('BoxPlot of Cloud9am')
plt.xlabel('Cloud9am')
plt.show()
```



```
In [61]: outlier_per(df, 'Cloud9am')
```

```
Out[61]: 0.0
```

```
In [62]: mean_cloud9am = df['Cloud9am'].mean()  
print(mean_cloud9am)
```

```
4.4474612602152455
```

```
In [63]: df['Cloud9am'].fillna(value = mean_cloud9am, inplace = True)
```

```
In [64]: df.isnull().sum()
```

```
Out[64]: Date          0
         Location      0
         MinTemp      1485
         MaxTemp      1261
         Rainfall      0
         Evaporation  62790
         Sunshine     69835
         WindGustDir   10326
         WindGustSpeed  0
         WindDir9am    10566
         WindDir3pm    4228
         WindSpeed9am  1767
         WindSpeed3pm  3062
         Humidity9am   0
         Humidity3pm   4507
         Pressure9am   0
         Pressure3pm  15028
         Cloud9am      0
         Cloud3pm     59358
         Temp9am       1767
         Temp3pm       3609
         RainToday     0
         RainTomorrow  0
         dtype: int64
```

Imputing missing values in Temp9am column using Median

```
In [65]: df[['Temp9am']]
```

Out[65]:

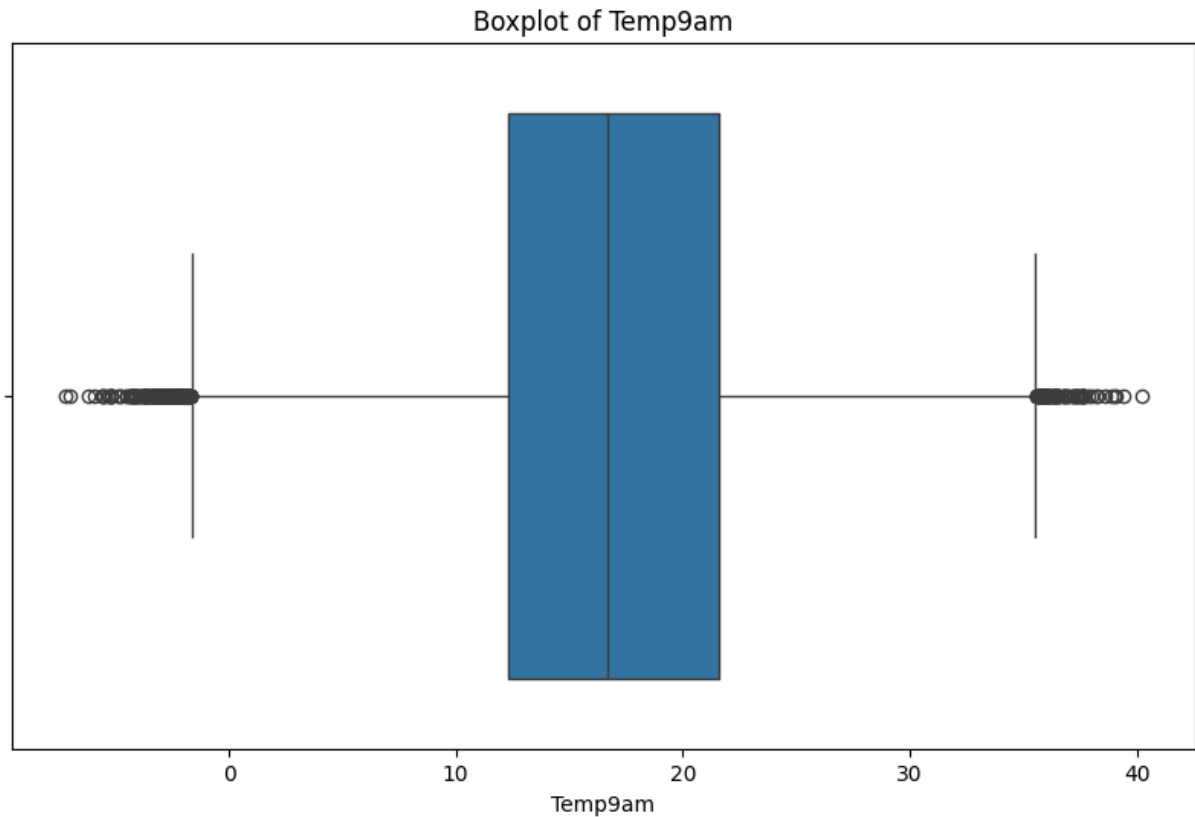
	Temp9am
0	16.9
1	17.2
2	21.0
3	18.1
4	17.8
...	...
145455	10.1
145456	10.9
145457	12.5
145458	15.1
145459	15.0

145460 rows × 1 columns

```
In [66]: df['Temp9am'].isnull().sum()
```

Out[66]: 1767

```
In [67]: plt.figure(figsize=(10, 6))
sns.boxplot(x=df['Temp9am'])
plt.title('Boxplot of Temp9am')
plt.xlabel('Temp9am')
plt.show()
```



```
In [68]: outlier_per(df, 'Temp9am')
```

Out[68]: 0.18011824556579129

```
In [69]: median_temp = df['Temp9am'].median()
print(median_temp)
```

16.7

```
In [70]: df['Temp9am'].fillna(value = median_temp , inplace = True)
```

```
In [71]: df.isnull().sum()
```

```
Out[71]: Date          0
        Location       0
        MinTemp        1485
        MaxTemp        1261
        Rainfall       0
        Evaporation    62790
        Sunshine       69835
        WindGustDir     10326
        WindGustSpeed   0
        WindDir9am     10566
        WindDir3pm     4228
        WindSpeed9am    1767
        WindSpeed3pm    3062
        Humidity9am     0
        Humidity3pm    4507
        Pressure9am     0
        Pressure3pm    15028
        Cloud9am       0
        Cloud3pm       59358
        Temp9am        0
        Temp3pm        3609
        RainToday      0
        RainTomorrow   0
        dtype: int64
```

## Imputing missing values in MinTemp column using Median

```
In [72]: df[['MinTemp']]
```

Out[72]:

MinTemp	
0	13.4
1	7.4
2	12.9
3	9.2
4	17.5
...	...
145455	2.8
145456	3.6
145457	5.4
145458	7.8
145459	14.9

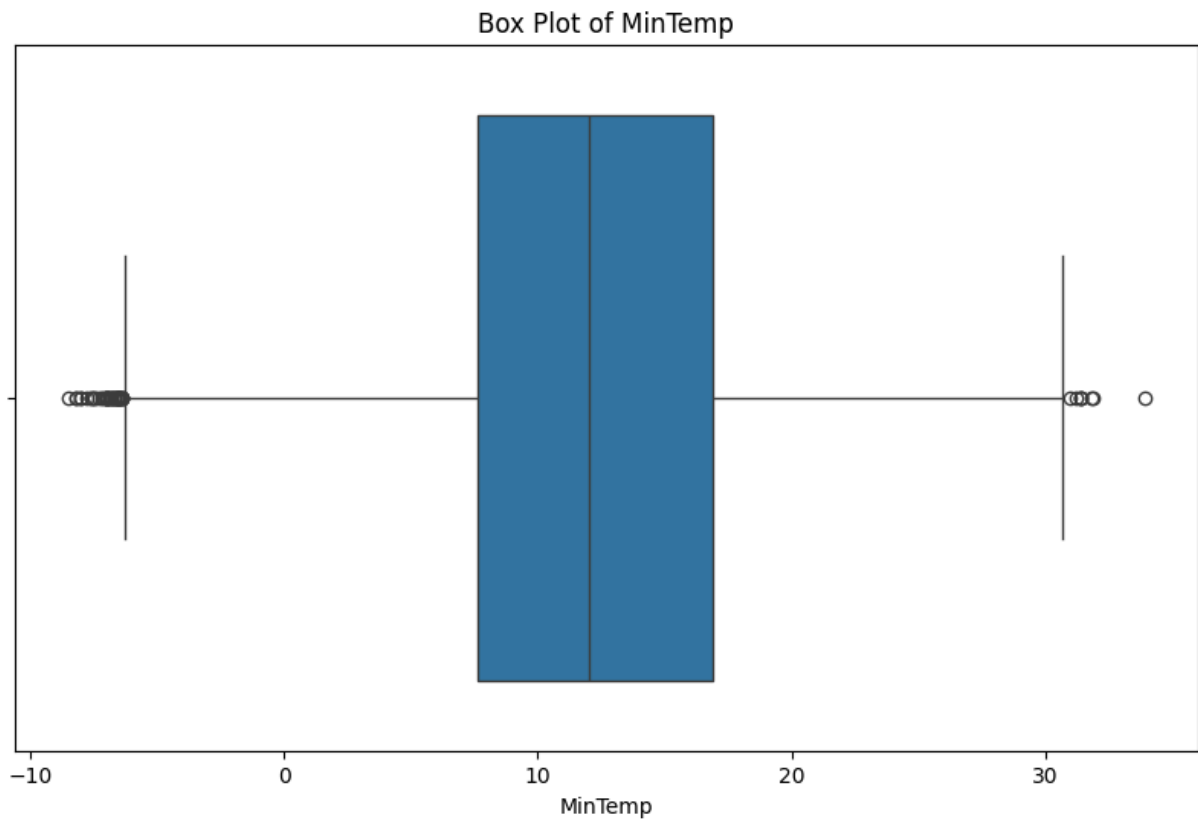
145460 rows × 1 columns

In [73]: `df['MinTemp'].isnull().sum()`

Out[73]: 1485

In [74]: 

```
plt.figure(figsize=(10,6))
sns.boxplot(x=df['MinTemp'])
plt.title('Box Plot of MinTemp')
plt.xlabel('MinTemp')
plt.show()
```



```
In [75]: outlier_per(df, 'MinTemp')
```

```
Out[75]: 0.037123607864705074
```

```
In [76]: median_mintemp = df['MinTemp'].median()  
print(median_mintemp)
```

```
12.0
```

```
In [77]: df['MinTemp'].fillna(value=median_mintemp, inplace= True)
```

```
In [78]: df.isnull().sum()
```



```
Out[78]: Date          0
         Location      0
         MinTemp       0
         MaxTemp      1261
         Rainfall      0
         Evaporation   62790
         Sunshine      69835
         WindGustDir    10326
         WindGustSpeed  0
         WindDir9am    10566
         WindDir3pm    4228
         WindSpeed9am   1767
         WindSpeed3pm   3062
         Humidity9am    0
         Humidity3pm    4507
         Pressure9am    0
         Pressure3pm   15028
         Cloud9am      0
         Cloud3pm      59358
         Temp9am       0
         Temp3pm       3609
         RainToday     0
         RainTomorrow  0
         dtype: int64
```

Imputing values for Maxtemp using Median

```
In [79]: df['MaxTemp']
```

```
Out[79]: 0          22.9
         1          25.1
         2          25.7
         3          28.0
         4          32.3
         ...
         145455      23.4
         145456      25.3
         145457      26.9
         145458      27.0
         145459      NaN
         Name: MaxTemp, Length: 145460, dtype: float64
```

```
In [80]: df['MaxTemp'].dtypes
```

```
Out[80]: dtype('float64')
```

```
In [81]: df['MaxTemp'].describe()
```

```
Out[81]: count    144199.000000  
        mean      23.221348  
        std       7.119049  
        min      -4.800000  
        25%      17.900000  
        50%      22.600000  
        75%      28.200000  
        max       48.100000  
        Name: MaxTemp, dtype: float64
```

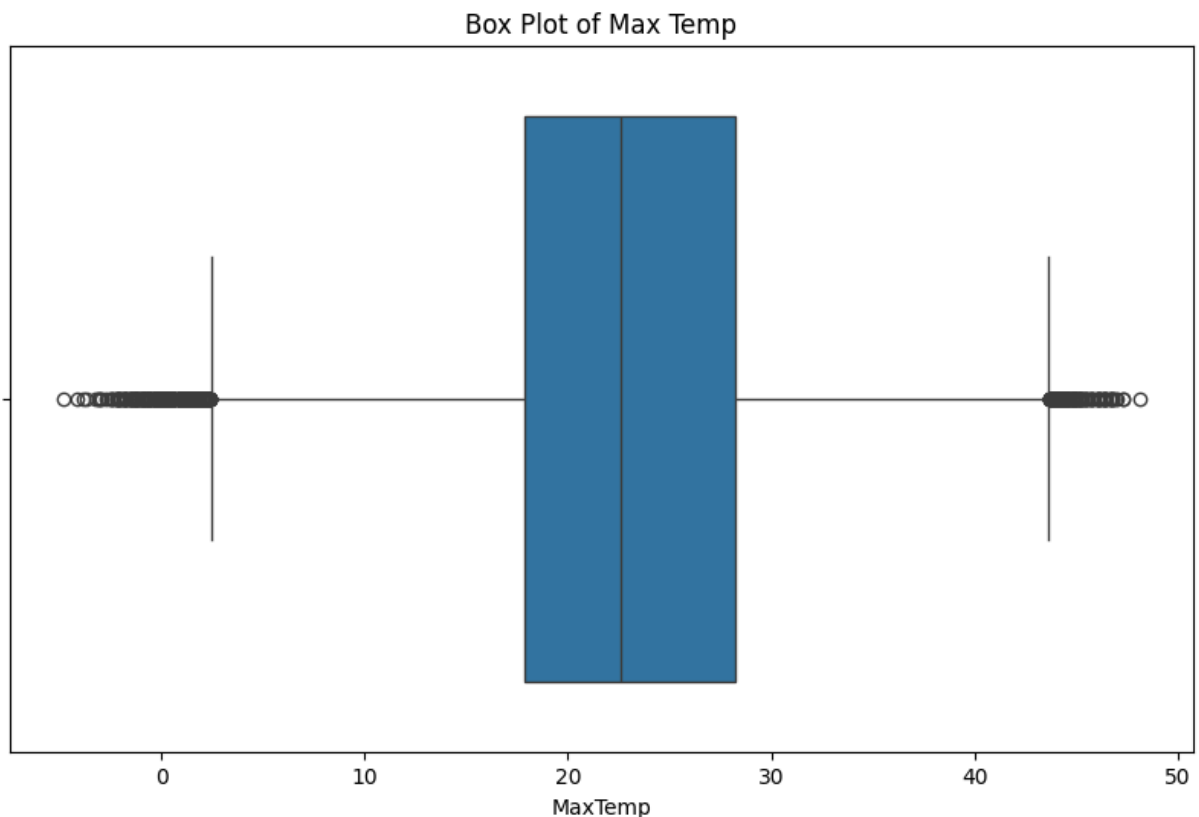
```
In [82]: df['MaxTemp'].isnull().sum()
```

```
Out[82]: 1261
```

```
In [83]: outlier_per(df, 'MaxTemp')
```

```
Out[83]: 0.3361748934414959
```

```
In [84]: plt.figure(figsize=(10,6))  
        sns.boxplot(x=df['MaxTemp'])  
        plt.title('Box Plot of Max Temp')  
        plt.xlabel('MaxTemp')  
        plt.show()
```



```
In [85]: median_maxtemp = df['MaxTemp'].median()  
        print('median_maxtemp')  
        df['MaxTemp'].fillna(value = median_maxtemp, inplace = True)
```

```
median_maxtemp
```

```
In [86]: df.isnull().sum()
```

```
Out[86]: Date                0
Location                  0
MinTemp                  0
MaxTemp                  0
Rainfall                 0
Evaporation             62790
Sunshine                 69835
WindGustDir             10326
WindGustSpeed            0
WindDir9am              10566
WindDir3pm               4228
WindSpeed9am            1767
WindSpeed3pm            3062
Humidity9am              0
Humidity3pm             4507
Pressure9am              0
Pressure3pm             15028
Cloud9am                 0
Cloud3pm                 59358
Temp9am                  0
Temp3pm                  3609
RainToday                0
RainTomorrow             0
dtype: int64
```

IMPUTING VALUES FOR Evaporation using Median

```
In [87]: df['Evaporation']
```

```
Out[87]: 0      NaN
1      NaN
2      NaN
3      NaN
4      NaN
..
145455 NaN
145456 NaN
145457 NaN
145458 NaN
145459 NaN
Name: Evaporation, Length: 145460, dtype: float64
```

```
In [88]: df['Evaporation'].dtypes
```

```
Out[88]: dtype('float64')
```

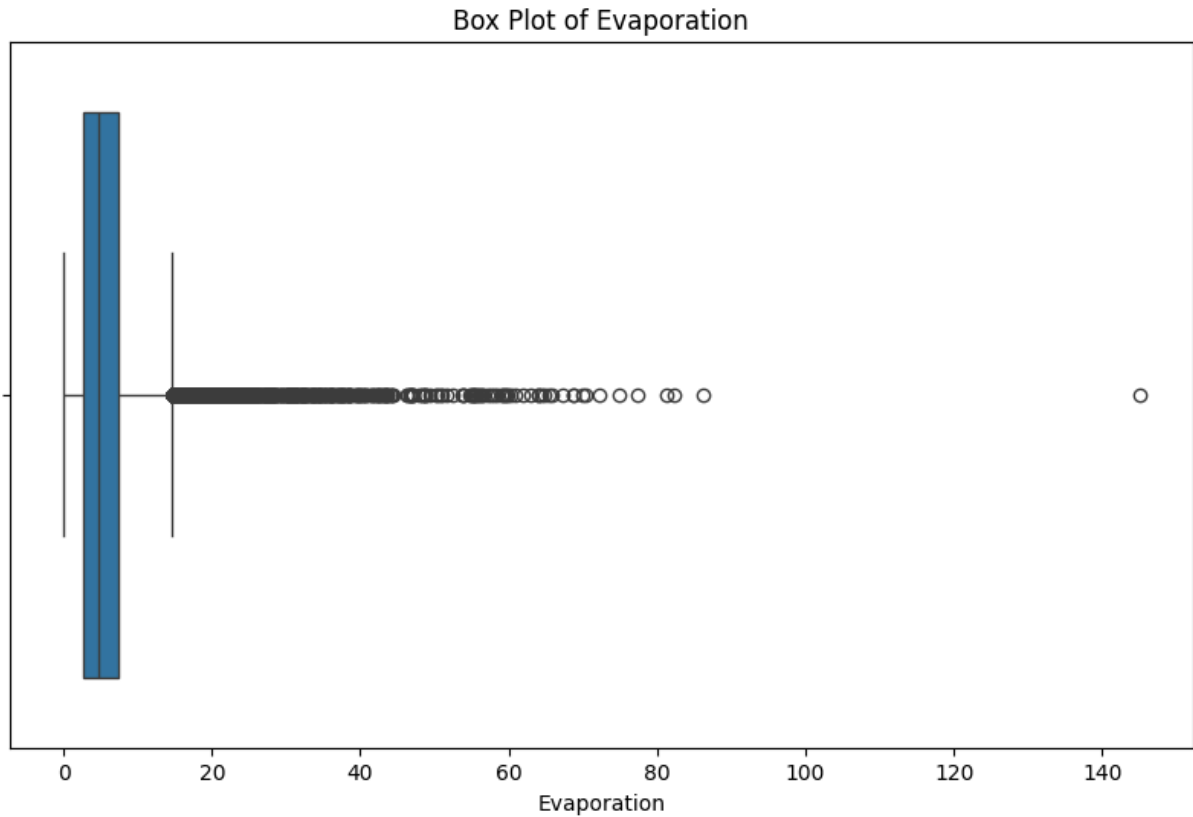
```
In [89]: df['Evaporation'].describe()
```

```
Out[89]: count    82670.000000
         mean      5.468232
         std       4.193704
         min       0.000000
         25%       2.600000
         50%       4.800000
         75%       7.400000
         max      145.000000
         Name: Evaporation, dtype: float64
```

```
In [90]: outlier_per(df, 'Evaporation')
```

```
Out[90]: 1.3715110683349374
```

```
In [91]: plt.figure(figsize=(10,6))
         sns.boxplot(x=df['Evaporation'])
         plt.title('Box Plot of Evaporation')
         plt.xlabel('Evaporation')
         plt.show()
```



```
In [92]: median_evaporation = df['Evaporation'].median()
```

```
In [93]: print(median_evaporation)
```

```
4.8
```

```
In [94]: df['Evaporation'].fillna(value = median_evaporation, inplace = True)
```

```
In [95]: df.isnull().sum()
```

```
Out[95]: Date          0
         Location      0
         MinTemp       0
         MaxTemp       0
         Rainfall      0
         Evaporation   0
         Sunshine      69835
         WindGustDir    10326
         WindGustSpeed  0
         WindDir9am     10566
         WindDir3pm     4228
         WindSpeed9am   1767
         WindSpeed3pm   3062
         Humidity9am    0
         Humidity3pm    4507
         Pressure9am    0
         Pressure3pm    15028
         Cloud9am       0
         Cloud3pm       59358
         Temp9am        0
         Temp3pm        3609
         RainToday      0
         RainTomorrow   0
         dtype: int64
```

```
In [96]: df[['Sunshine']]
```

```
Out[96]:
```

	Sunshine
0	NaN
1	NaN
2	NaN
3	NaN
4	NaN
...	...
145455	NaN
145456	NaN
145457	NaN
145458	NaN
145459	NaN

145460 rows × 1 columns

```
In [97]: df['Sunshine'].dtypes
```

```
Out[97]: dtype('float64')
```

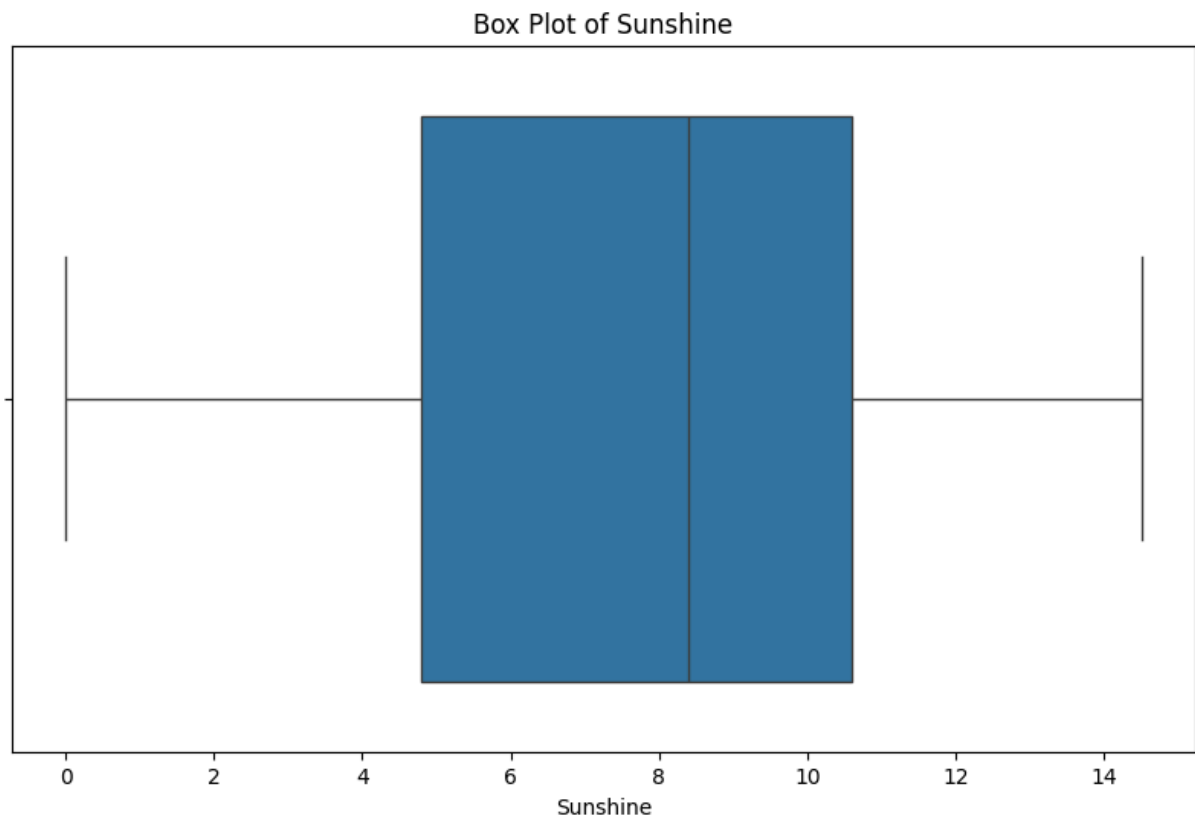
```
In [98]: df['Sunshine'].describe()
```

```
Out[98]: count    75625.000000  
mean         7.611178  
std          3.785483  
min           0.000000  
25%          4.800000  
50%          8.400000  
75%         10.600000  
max         14.500000  
Name: Sunshine, dtype: float64
```

```
In [99]: outlier_per(df, 'Sunshine')
```

```
Out[99]: 0.0
```

```
In [100]: plt.figure(figsize=(10,6))  
sns.boxplot(x=df['Sunshine'])  
plt.title('Box Plot of Sunshine')  
plt.xlabel('Sunshine')  
plt.show()
```



```
In [101]: mean_sunshine = df['Sunshine'].mean()
```

```
In [102]: df['Sunshine'].fillna(value = mean_sunshine, inplace = True)
```

```
In [103]: df.isnull().sum()
```

```
Out[103... Date          0
Location        0
MinTemp         0
MaxTemp         0
Rainfall        0
Evaporation     0
Sunshine        0
WindGustDir     10326
WindGustSpeed   0
WindDir9am     10566
WindDir3pm     4228
WindSpeed9am    1767
WindSpeed3pm    3062
Humidity9am     0
Humidity3pm    4507
Pressure9am     0
Pressure3pm    15028
Cloud9am        0
Cloud3pm       59358
Temp9am         0
Temp3pm        3609
RainToday       0
RainTomorrow    0
dtype: int64
```

```
In [104... df.describe(include = ['object'])
```

```
Out[104...
      Date  Location  WindGustDir  WindDir9am  WindDir3pm
count  145460   145460      135134      134894      141232
unique    3436        49          16          16          16
top  2013-11-12  Canberra          W          N          SE
freq      49      3436       9915      11758      10838
```

Imputing values for WindSpeed using median

```
In [105... df[['WindSpeed9am']]]
```

Out[105...

**WindSpeed9am**

<b>0</b>	20.0
<b>1</b>	4.0
<b>2</b>	19.0
<b>3</b>	11.0
<b>4</b>	7.0
...	...
<b>145455</b>	13.0
<b>145456</b>	13.0
<b>145457</b>	9.0
<b>145458</b>	13.0
<b>145459</b>	17.0

145460 rows × 1 columns

In [106...

`df[['WindSpeed9am']].dtypes`

Out[106...

WindSpeed9am    float64  
dtype: object

In [107...

`outlier_per(df, 'WindSpeed9am')`

Out[107...

1.249140657225354

In [108...

```
median_windspeed9am = df['WindSpeed9am'].median()
print(median_windspeed9am)
```

13.0

In [109...

`df['WindSpeed9am'].fillna(value = median_windspeed9am, inplace = True)`

In [110...

`df.isnull().sum()`



```
Out[110... Date          0
Location        0
MinTemp         0
MaxTemp         0
Rainfall        0
Evaporation     0
Sunshine        0
WindGustDir     10326
WindGustSpeed    0
WindDir9am      10566
WindDir3pm      4228
WindSpeed9am     0
WindSpeed3pm    3062
Humidity9am      0
Humidity3pm     4507
Pressure9am      0
Pressure3pm     15028
Cloud9am         0
Cloud3pm        59358
Temp9am          0
Temp3pm         3609
RainToday        0
RainTomorrow     0
dtype: int64
```

Imputing values for WindSpeed3pm using mean

```
In [111... df[['WindSpeed3pm']]]
```

Out[111...

WindSpeed3pm	
0	24.0
1	22.0
2	26.0
3	9.0
4	20.0
...	...
145455	11.0
145456	9.0
145457	9.0
145458	7.0
145459	17.0

145460 rows × 1 columns

```
In [112... df[['WindSpeed3pm']].describe()
```

Out[112...

**WindSpeed3pm**

<b>count</b>	142398.000000
<b>mean</b>	18.662657
<b>std</b>	8.809800
<b>min</b>	0.000000
<b>25%</b>	13.000000
<b>50%</b>	19.000000
<b>75%</b>	24.000000
<b>max</b>	87.000000

In [113... `df[['WindSpeed3pm']].dtypes`Out[113... WindSpeed3pm float64  
dtype: objectIn [114... `outlier_per(df, 'WindSpeed3pm')`

Out[114... 1.734497456345387

In [115... `mean_windspeed3pm = df['WindSpeed3pm'].mean()`In [116... `df['WindSpeed3pm'].fillna(value = mean_windspeed3pm, inplace = True)`In [117... `df.isnull().sum()`

```
Out[117... Date          0
Location        0
MinTemp         0
MaxTemp         0
Rainfall        0
Evaporation     0
Sunshine        0
WindGustDir     10326
WindGustSpeed   0
WindDir9am      10566
WindDir3pm      4228
WindSpeed9am    0
WindSpeed3pm    0
Humidity9am     0
Humidity3pm     4507
Pressure9am     0
Pressure3pm     15028
Cloud9am        0
Cloud3pm        59358
Temp9am         0
Temp3pm         3609
RainToday       0
RainTomorrow    0
dtype: int64
```

Imputing Values for Humidity3pm using Mean

```
In [118... df[['Humidity3pm']]]
```

Out[118...

Humidity3pm	
0	22.0
1	25.0
2	30.0
3	16.0
4	33.0
...	...
145455	24.0
145456	21.0
145457	24.0
145458	24.0
145459	36.0

145460 rows × 1 columns

```
In [119... df['Humidity3pm'].isnull().sum()
```

Out[119... 4507

```
In [120... df['Humidity3pm'].describe()
```

```
Out[120... count    140953.000000
mean       51.539116
std        20.795902
min         0.000000
25%        37.000000
50%        52.000000
75%        66.000000
max       100.000000
Name: Humidity3pm, dtype: float64
```

```
In [121... df['Humidity3pm'].dtypes
```

```
Out[121... dtype('float64')
```

```
In [122... outlier_per(df, 'Humidity3pm')
```

```
Out[122... 0.0
```

```
In [123... mean_humidity3pm = df['Humidity3pm'].mean()
```

```
In [124... df['Humidity3pm'].fillna(value = mean_humidity3pm, inplace = True)
```

```
In [125... df.isnull().sum()
```

```
Out[125... Date            0
Location          0
MinTemp           0
MaxTemp           0
Rainfall          0
Evaporation       0
Sunshine          0
WindGustDir      10326
WindGustSpeed     0
WindDir9am       10566
WindDir3pm       4228
WindSpeed9am     0
WindSpeed3pm     0
Humidity9am      0
Humidity3pm      0
Pressure9am      0
Pressure3pm     15028
Cloud9am         0
Cloud3pm        59358
Temp9am          0
Temp3pm         3609
RainToday        0
RainTomorrow     0
dtype: int64
```

Imputing values for Pressure3pm using median

```
In [126... df['Pressure3pm']
```

```
Out[126... 0          1007.1
1          1007.8
2          1008.7
3          1012.8
4          1006.0
...
145455     1020.3
145456     1019.1
145457     1016.8
145458     1016.5
145459     1017.9
Name: Pressure3pm, Length: 145460, dtype: float64
```

```
In [127... df['Pressure3pm'].describe()
```

```
Out[127... count    130432.000000
mean      1015.255889
std        7.037414
min        977.100000
25%       1010.400000
50%       1015.200000
75%       1020.000000
max       1039.600000
Name: Pressure3pm, dtype: float64
```

```
In [128... df['Pressure3pm'].dtypes
```

```
Out[128... dtype('float64')
```

```
In [129... outlier_per(df, 'Pressure3pm')
```

```
Out[129... 0.6317888079197029
```

```
In [130... median_pressure3pm = df['Pressure3pm'].median()
```

```
In [131... print(median_pressure3pm)
```

```
1015.2
```

```
In [132... df['Pressure3pm'].fillna(value = median_pressure3pm, inplace = True)
```

```
In [133... df.isnull().sum()
```

```
Out[133... Date          0
Location        0
MinTemp         0
MaxTemp         0
Rainfall        0
Evaporation     0
Sunshine        0
WindGustDir     10326
WindGustSpeed   0
WindDir9am     10566
WindDir3pm     4228
WindSpeed9am    0
WindSpeed3pm    0
Humidity9am     0
Humidity3pm     0
Pressure9am     0
Pressure3pm     0
Cloud9am        0
Cloud3pm       59358
Temp9am         0
Temp3pm        3609
RainToday       0
RainTomorrow    0
dtype: int64
```

Imputing values for Cloud3pm using Mean

```
In [134... df[['Cloud3pm']]]
```

Out[134...

Cloud3pm	
0	NaN
1	NaN
2	2.0
3	NaN
4	8.0
...	...
145455	NaN
145456	NaN
145457	NaN
145458	2.0
145459	8.0

145460 rows × 1 columns

```
In [135... df['Cloud3pm'].describe()
```

```
Out[135... count      86102.000000
          mean        4.509930
          std         2.720357
          min         0.000000
          25%         2.000000
          50%         5.000000
          75%         7.000000
          max         9.000000
          Name: Cloud3pm, dtype: float64
```

```
In [136... df['Cloud3pm'].dtypes
```

```
Out[136... dtype('float64')
```

```
In [137... outlier_per(df, 'Cloud3pm')
```

```
Out[137... 0.0
```

```
In [138... mean_cloud3pm = df['Cloud3pm'].mean()
```

```
In [139... df['Cloud3pm'].fillna(value = mean_cloud3pm, inplace = True)
```

```
In [140... df.isnull().sum()
```

```
Out[140... Date          0
          Location      0
          MinTemp       0
          MaxTemp       0
          Rainfall      0
          Evaporation   0
          Sunshine      0
          WindGustDir    10326
          WindGustSpeed  0
          WindDir9am     10566
          WindDir3pm     4228
          WindSpeed9am   0
          WindSpeed3pm   0
          Humidity9am    0
          Humidity3pm    0
          Pressure9am    0
          Pressure3pm    0
          Cloud9am       0
          Cloud3pm       0
          Temp9am        0
          Temp3pm       3609
          RainToday      0
          RainTomorrow   0
          dtype: int64
```

Imputing values for Temp3pm using Mean

```
In [141... df[['Temp3pm']]]
```

Out[141...

Temp3pm	
0	21.8
1	24.3
2	23.2
3	26.5
4	29.7
...	...
145455	22.4
145456	24.5
145457	26.1
145458	26.0
145459	20.9

145460 rows × 1 columns

In [142...

df['Temp3pm'].describe()

Out[142...

```
count    141851.00000
mean      21.68339
std        6.93665
min       -5.40000
25%       16.60000
50%       21.10000
75%       26.40000
max       46.70000
Name: Temp3pm, dtype: float64
```

In [143...

df['Temp3pm'].dtypes

Out[143...

dtype('float64')

In [144...

outlier\_per(df, 'Temp3pm')

Out[144...

0.525230303863605

In [145...

mean\_temp3pm = df['Temp3pm'].mean()

In [146...

print(mean\_temp3pm)

21.68339031800974

In [147...

df['Temp3pm'].fillna(value = mean\_temp3pm, inplace = True)

In [148...

df.isnull().sum()



```
Out[148... Date          0
Location        0
MinTemp         0
MaxTemp         0
Rainfall        0
Evaporation     0
Sunshine        0
WindGustDir     10326
WindGustSpeed   0
WindDir9am      10566
WindDir3pm      4228
WindSpeed9am    0
WindSpeed3pm    0
Humidity9am     0
Humidity3pm     0
Pressure9am     0
Pressure3pm     0
Cloud9am        0
Cloud3pm        0
Temp9am         0
Temp3pm         0
RainToday       0
RainTomorrow    0
dtype: int64
```

Imputing values for WindGustDir using MMode

```
In [149... df[['WindGustDir']]]
```

Out[149...

WindGustDir	
0	W
1	WNW
2	WSW
3	NE
4	W
...	...
145455	E
145456	NNW
145457	N
145458	SE
145459	NaN

145460 rows × 1 columns

```
In [150... df['WindGustDir'].value_counts()
```

```
Out[150...] WindGustDir
           W      9915
           SE     9418
           N      9313
           SSE    9216
           E      9181
           S      9168
           WSW    9069
           SW     8967
           SSW    8736
           WNW    8252
           NW     8122
           ENE    8104
           ESE    7372
           NE     7133
           NNW    6620
           NNE    6548
           Name: count, dtype: int64
```

```
In [151...] df['WindGustDir'].describe(include=['object'])
```

```
Out[151...] count      135134
           unique         16
           top            W
           freq          9915
           Name: WindGustDir, dtype: object
```

```
In [152...] df['WindGustDir'].dtypes
```

```
Out[152...] dtype('O')
```

```
In [153...] df['WindGustDir'].isnull().sum()
```

```
Out[153...] 10326
```

```
In [154...] mode_windgustdir = df['WindGustDir'].mode()[0]
```

```
In [155...] print(mode_windgustdir)
```

W

```
In [156...] df['WindGustDir'].fillna(value = mode_windgustdir, inplace = True)
```

```
In [157...] df.isnull().sum()
```

Out[157...      Date                      0  
                 Location                0  
                 MinTemp                0  
                 MaxTemp                0  
                 Rainfall                0  
                 Evaporation            0  
                 Sunshine                0  
                 WindGustDir            0  
                 WindGustSpeed        0  
                 WindDir9am        10566  
                 WindDir3pm        4228  
                 WindSpeed9am        0  
                 WindSpeed3pm        0  
                 Humidity9am        0  
                 Humidity3pm        0  
                 Pressure9am        0  
                 Pressure3pm        0  
                 Cloud9am            0  
                 Cloud3pm            0  
                 Temp9am            0  
                 Temp3pm            0  
                 RainToday            0  
                 RainTomorrow        0  
                 dtype: int64

Imputing Values For WindDir9am using mode

```
In [158... df[['WindDir9am']]]
```

Out[158...

WindDir9am	
0	W
1	NNW
2	W
3	SE
4	ENE
...	...
145455	SE
145456	SE
145457	SE
145458	SSE
145459	ESE

145460 rows × 1 columns

```
In [159... df[['WindDir9am']].value_counts()
```

```
Out[159...] WindDir9am
           N          11758
           SE          9287
           E          9176
           SSE         9112
           NW          8749
           S           8659
           W           8459
           SW          8423
           NNE         8129
           NNW         7980
           ENE         7836
           NE          7671
           ESE         7630
           SSW         7587
           WNW         7414
           WSW         7024
           Name: count, dtype: int64
```

```
In [160...] df[['WindDir9am']].dtypes
```

```
Out[160...] WindDir9am    object
           dtype: object
```

```
In [161...] mode_WindDir9am = df['WindDir9am'].mode()[0]
```

```
In [162...] print(mode_WindDir9am)
```

N

```
In [163...] df['WindDir9am'].fillna(value = mode_WindDir9am, inplace = True)
```

```
In [164...] df.isnull().sum()
```

Out[164...    Date                    0  
              Location           0  
              MinTemp            0  
              MaxTemp            0  
              Rainfall            0  
              Evaporation        0  
              Sunshine            0  
              WindGustDir        0  
              WindGustSpeed      0  
              WindDir9am        0  
              WindDir3pm        4228  
              WindSpeed9am       0  
              WindSpeed3pm       0  
              Humidity9am        0  
              Humidity3pm        0  
              Pressure9am        0  
              Pressure3pm        0  
              Cloud9am           0  
              Cloud3pm           0  
              Temp9am            0  
              Temp3pm            0  
              RainToday          0  
              RainTomorrow       0  
              dtype: int64

Imputing Values for WindDir3pm using mode

In [165...    `df[['WindDir3pm']]`

Out[165...    **WindDir3pm**

<b>0</b>	WNW
<b>1</b>	WSW
<b>2</b>	WSW
<b>3</b>	E
<b>4</b>	NW
...	...
<b>145455</b>	ENE
<b>145456</b>	N
<b>145457</b>	WNW
<b>145458</b>	N
<b>145459</b>	ESE

145460 rows × 1 columns

In [166...    `df[['WindDir3pm']].value_counts()`

```
Out[166... WindDir3pm
           SE      10838
           W      10110
           S      9926
           WSW     9518
           SSE     9399
           SW      9354
           N      8890
           WNW     8874
           NW      8610
           ESE     8505
           E      8472
           NE      8263
           SSW     8156
           NNW     7870
           ENE     7857
           NNE     6590
Name: count, dtype: int64
```

```
In [167... df['WindDir3pm'].dtypes
```

```
Out[167... dtype('O')
```

```
In [168... mode_WindDir3pm = df['WindDir3pm'].mode()[0]
```

```
In [169... print(mode_WindDir3pm)
```

SE

```
In [170... df['WindDir3pm'].fillna(value = mode_WindDir3pm, inplace = True)
```

```
In [171... df.isnull().sum()
```

```
Out[171... Date          0
          Location      0
          MinTemp       0
          MaxTemp       0
          Rainfall      0
          Evaporation    0
          Sunshine      0
          WindGustDir    0
          WindGustSpeed  0
          WindDir9am     0
          WindDir3pm     0
          WindSpeed9am   0
          WindSpeed3pm   0
          Humidity9am    0
          Humidity3pm    0
          Pressure9am    0
          Pressure3pm    0
          Cloud9am       0
          Cloud3pm       0
          Temp9am        0
          Temp3pm        0
          RainToday      0
          RainTomorrow   0
          dtype: int64
```

Data Type Conversion: Convert date and categorical variables to appropriate formats. The date might be segmented into year, month, and day components for more detailed analysis

```
In [172... df['Date'] = pd.to_datetime(df['Date'])
df['Year'] = df['Date'].dt.year
df['Month'] = df['Date'].dt.month
df['Day'] = df['Date'].dt.day
```

```
In [173... df[['Date']]]
```

Out[173...

	Date
0	2008-12-01
1	2008-12-02
2	2008-12-03
3	2008-12-04
4	2008-12-05
...	...
145455	2017-06-21
145456	2017-06-22
145457	2017-06-23
145458	2017-06-24
145459	2017-06-25

145460 rows × 1 columns

In [174...


```
df
```



Out[174...

	Date	Location	MinTemp	MaxTemp	Rainfall	Evaporation	Sunshine	WindGust
<b>0</b>	2008-12-01	Albury	13.4	22.9	0.6	4.8	7.611178	
<b>1</b>	2008-12-02	Albury	7.4	25.1	0.0	4.8	7.611178	WN
<b>2</b>	2008-12-03	Albury	12.9	25.7	0.0	4.8	7.611178	WS
<b>3</b>	2008-12-04	Albury	9.2	28.0	0.0	4.8	7.611178	
<b>4</b>	2008-12-05	Albury	17.5	32.3	1.0	4.8	7.611178	
...	...	...	...	...	...	...	...	
<b>145455</b>	2017-06-21	Uluru	2.8	23.4	0.0	4.8	7.611178	
<b>145456</b>	2017-06-22	Uluru	3.6	25.3	0.0	4.8	7.611178	NN
<b>145457</b>	2017-06-23	Uluru	5.4	26.9	0.0	4.8	7.611178	
<b>145458</b>	2017-06-24	Uluru	7.8	27.0	0.0	4.8	7.611178	
<b>145459</b>	2017-06-25	Uluru	14.9	22.6	0.0	4.8	7.611178	

145460 rows × 26 columns



In [175...

df['Location'].unique()

Out[175...

```
array(['Albury', 'BadgerysCreek', 'Cobar', 'CoffsHarbour', 'Moree',
      'Newcastle', 'NorahHead', 'NorfolkIsland', 'Penrith', 'Richmond',
      'Sydney', 'SydneyAirport', 'WaggaWagga', 'Williamtown',
      'Wollongong', 'Canberra', 'Tuggeranong', 'MountGinini', 'Ballarat',
      'Bendigo', 'Sale', 'MelbourneAirport', 'Melbourne', 'Mildura',
      'Nhil', 'Portland', 'Watsonia', 'Dartmoor', 'Brisbane', 'Cairns',
      'GoldCoast', 'Townsville', 'Adelaide', 'MountGambier', 'Nuriootpa',
      'Woomera', 'Albany', 'Witchcliffe', 'PearceRAAF', 'PerthAirport',
      'Perth', 'SalmonGums', 'Walpole', 'Hobart', 'Launceston',
      'AliceSprings', 'Darwin', 'Katherine', 'Uluru'], dtype=object)
```

In [176...

```
locations = ['Albury', 'BadgerysCreek', 'Cobar', 'CoffsHarbour', 'Moree',
             'Newcastle', 'NorahHead', 'NorfolkIsland', 'Penrith', 'Richmond',
             'Sydney', 'SydneyAirport', 'WaggaWagga', 'Williamtown',
             'Wollongong', 'Canberra', 'Tuggeranong', 'MountGinini', 'Ballarat',
             'Bendigo', 'Sale', 'MelbourneAirport', 'Melbourne', 'Mildura',
             'Nhil', 'Portland', 'Watsonia', 'Dartmoor', 'Brisbane', 'Cairns',
```

```
'GoldCoast', 'Townsville', 'Adelaide', 'MountGambier', 'Nuriootpa',  
'Woomera', 'Albany', 'Witchcliffe', 'PearceRAAF', 'PerthAirport',  
'Perth', 'SalmonGums', 'Walpole', 'Hobart', 'Launceston',  
'AliceSprings', 'Darwin', 'Katherine', 'Uluru']  
  
location_dict = {location: index + 1 for index, location in enumerate(locations)}  
  
df['Location'] = df['Location'].map(location_dict)  
  
print(df)
```

	Date	Location	MinTemp	MaxTemp	Rainfall	Evaporation	\
0	2008-12-01	1	13.4	22.9	0.6	4.8	
1	2008-12-02	1	7.4	25.1	0.0	4.8	
2	2008-12-03	1	12.9	25.7	0.0	4.8	
3	2008-12-04	1	9.2	28.0	0.0	4.8	
4	2008-12-05	1	17.5	32.3	1.0	4.8	
...	...	...	...	...	...	...	
145455	2017-06-21	49	2.8	23.4	0.0	4.8	
145456	2017-06-22	49	3.6	25.3	0.0	4.8	
145457	2017-06-23	49	5.4	26.9	0.0	4.8	
145458	2017-06-24	49	7.8	27.0	0.0	4.8	
145459	2017-06-25	49	14.9	22.6	0.0	4.8	

	Sunshine	WindGustDir	WindGustSpeed	WindDir9am	...	Pressure3pm	\
0	7.611178	W	44.0	W	...	1007.1	
1	7.611178	WNW	44.0	NNW	...	1007.8	
2	7.611178	WSW	46.0	W	...	1008.7	
3	7.611178	NE	24.0	SE	...	1012.8	
4	7.611178	W	41.0	ENE	...	1006.0	
...	...	...	...	...	...	...	
145455	7.611178	E	31.0	SE	...	1020.3	
145456	7.611178	NNW	22.0	SE	...	1019.1	
145457	7.611178	N	37.0	SE	...	1016.8	
145458	7.611178	SE	28.0	SSE	...	1016.5	
145459	7.611178	W	39.0	ESE	...	1017.9	

	Cloud9am	Cloud3pm	Temp9am	Temp3pm	RainToday	RainTomorrow	Year	\
0	8.000000	4.50993	16.9	21.8	0	0	2008	
1	4.447461	4.50993	17.2	24.3	0	0	2008	
2	4.447461	2.000000	21.0	23.2	0	0	2008	
3	4.447461	4.50993	18.1	26.5	0	0	2008	
4	7.000000	8.000000	17.8	29.7	0	0	2008	
...	...	...	...	...	...	...	...	
145455	4.447461	4.50993	10.1	22.4	0	0	2017	
145456	4.447461	4.50993	10.9	24.5	0	0	2017	
145457	4.447461	4.50993	12.5	26.1	0	0	2017	
145458	3.000000	2.000000	15.1	26.0	0	0	2017	
145459	8.000000	8.000000	15.0	20.9	0	0	2017	

	Month	Day
0	12	1
1	12	2
2	12	3
3	12	4
4	12	5
...	...	...
145455	6	21
145456	6	22
145457	6	23
145458	6	24
145459	6	25

[145460 rows x 26 columns]

In [177... df['WindGustDir'].unique()

```
Out[177... array(['W', 'WNW', 'WSW', 'NE', 'NNW', 'N', 'NNE', 'SW', 'ENE', 'SSE',  
      'S', 'NW', 'SE', 'ESE', 'E', 'SSW'], dtype=object)
```

```
In [178... WGD = ['W', 'WNW', 'WSW', 'NE', 'NNW', 'N', 'NNE', 'SW', 'ENE', 'SSE',  
      'S', 'NW', 'SE', 'ESE', 'E', 'SSW']  
  
wgd_dict = {WGD: index + 1 for index, WGD in enumerate(WGD )}  
  
df['WindGustDir'] = df['WindGustDir'].map(wgd_dict)  
  
print(df)
```

	Date	Location	MinTemp	MaxTemp	Rainfall	Evaporation	\
0	2008-12-01	1	13.4	22.9	0.6	4.8	
1	2008-12-02	1	7.4	25.1	0.0	4.8	
2	2008-12-03	1	12.9	25.7	0.0	4.8	
3	2008-12-04	1	9.2	28.0	0.0	4.8	
4	2008-12-05	1	17.5	32.3	1.0	4.8	
...	...	...	...	...	...	...	
145455	2017-06-21	49	2.8	23.4	0.0	4.8	
145456	2017-06-22	49	3.6	25.3	0.0	4.8	
145457	2017-06-23	49	5.4	26.9	0.0	4.8	
145458	2017-06-24	49	7.8	27.0	0.0	4.8	
145459	2017-06-25	49	14.9	22.6	0.0	4.8	

	Sunshine	WindGustDir	WindGustSpeed	WindDir9am	... Pressure3pm	\
0	7.611178	1	44.0	W	1007.1	
1	7.611178	2	44.0	NNW	1007.8	
2	7.611178	3	46.0	W	1008.7	
3	7.611178	4	24.0	SE	1012.8	
4	7.611178	1	41.0	ENE	1006.0	
...	...	...	...	...	...	
145455	7.611178	15	31.0	SE	1020.3	
145456	7.611178	5	22.0	SE	1019.1	
145457	7.611178	6	37.0	SE	1016.8	
145458	7.611178	13	28.0	SSE	1016.5	
145459	7.611178	1	39.0	ESE	1017.9	

	Cloud9am	Cloud3pm	Temp9am	Temp3pm	RainToday	RainTomorrow	Year	\
0	8.000000	4.50993	16.9	21.8	0	0	2008	
1	4.447461	4.50993	17.2	24.3	0	0	2008	
2	4.447461	2.00000	21.0	23.2	0	0	2008	
3	4.447461	4.50993	18.1	26.5	0	0	2008	
4	7.000000	8.00000	17.8	29.7	0	0	2008	
...	...	...	...	...	...	...	...	
145455	4.447461	4.50993	10.1	22.4	0	0	2017	
145456	4.447461	4.50993	10.9	24.5	0	0	2017	
145457	4.447461	4.50993	12.5	26.1	0	0	2017	
145458	3.000000	2.00000	15.1	26.0	0	0	2017	
145459	8.000000	8.00000	15.0	20.9	0	0	2017	

	Month	Day
0	12	1
1	12	2
2	12	3
3	12	4
4	12	5
...	...	...
145455	6	21
145456	6	22
145457	6	23
145458	6	24
145459	6	25

[145460 rows x 26 columns]

In [179... df['WindDir9am'].unique()

```
Out[179... array(['W', 'NNW', 'SE', 'ENE', 'SW', 'SSE', 'S', 'NE', 'N', 'SSW', 'WSW',  
      'ESE', 'E', 'NW', 'WNW', 'NNE'], dtype=object)
```

```
In [180... WGD9am = ['W', 'NNW', 'SE', 'ENE', 'SW', 'SSE', 'S', 'NE', 'N', 'SSW', 'WSW',  
          'ESE', 'E', 'NW', 'WNW', 'NNE']  
  
wgd9am_dict = {WGD9am: index + 1 for index, WGD9am in enumerate(WGD9am )}  
  
df['WindDir9am'] = df['WindDir9am'].map(wgd9am_dict)  
  
print(df)
```

	Date	Location	MinTemp	MaxTemp	Rainfall	Evaporation	\
0	2008-12-01	1	13.4	22.9	0.6	4.8	
1	2008-12-02	1	7.4	25.1	0.0	4.8	
2	2008-12-03	1	12.9	25.7	0.0	4.8	
3	2008-12-04	1	9.2	28.0	0.0	4.8	
4	2008-12-05	1	17.5	32.3	1.0	4.8	
...	...	...	...	...	...	...	
145455	2017-06-21	49	2.8	23.4	0.0	4.8	
145456	2017-06-22	49	3.6	25.3	0.0	4.8	
145457	2017-06-23	49	5.4	26.9	0.0	4.8	
145458	2017-06-24	49	7.8	27.0	0.0	4.8	
145459	2017-06-25	49	14.9	22.6	0.0	4.8	

	Sunshine	WindGustDir	WindGustSpeed	WindDir9am	... Pressure3pm	\
0	7.611178	1	44.0	1	1007.1	
1	7.611178	2	44.0	2	1007.8	
2	7.611178	3	46.0	1	1008.7	
3	7.611178	4	24.0	3	1012.8	
4	7.611178	1	41.0	4	1006.0	
...	...	...	...	...	...	
145455	7.611178	15	31.0	3	1020.3	
145456	7.611178	5	22.0	3	1019.1	
145457	7.611178	6	37.0	3	1016.8	
145458	7.611178	13	28.0	6	1016.5	
145459	7.611178	1	39.0	12	1017.9	

	Cloud9am	Cloud3pm	Temp9am	Temp3pm	RainToday	RainTomorrow	Year	\
0	8.000000	4.50993	16.9	21.8	0	0	2008	
1	4.447461	4.50993	17.2	24.3	0	0	2008	
2	4.447461	2.00000	21.0	23.2	0	0	2008	
3	4.447461	4.50993	18.1	26.5	0	0	2008	
4	7.000000	8.00000	17.8	29.7	0	0	2008	
...	...	...	...	...	...	...	...	
145455	4.447461	4.50993	10.1	22.4	0	0	2017	
145456	4.447461	4.50993	10.9	24.5	0	0	2017	
145457	4.447461	4.50993	12.5	26.1	0	0	2017	
145458	3.000000	2.00000	15.1	26.0	0	0	2017	
145459	8.000000	8.00000	15.0	20.9	0	0	2017	

	Month	Day
0	12	1
1	12	2
2	12	3
3	12	4
4	12	5
...	...	...
145455	6	21
145456	6	22
145457	6	23
145458	6	24
145459	6	25

[145460 rows x 26 columns]

In [181... df['WindDir3pm'].unique()

```
Out[181...] array(['WNW', 'WSW', 'E', 'NW', 'W', 'SSE', 'ESE', 'ENE', 'NNW', 'SSW',  
      'SW', 'SE', 'N', 'S', 'NNE', 'NE'], dtype=object)
```

```
In [182...] WGD3pm = ['WNW', 'WSW', 'E', 'NW', 'W', 'SSE', 'ESE', 'ENE', 'NNW', 'SSW',  
          'SW', 'SE', 'N', 'S', 'NNE', 'NE']  
  
wgd3pm_dict = {WGD3pm: index + 1 for index, WGD3pm in enumerate(WGD3pm )}  
  
df['WindDir3pm'] = df['WindDir3pm'].map(wgd3pm_dict)  
  
print(df)
```



	Date	Location	MinTemp	MaxTemp	Rainfall	Evaporation	\
0	2008-12-01	1	13.4	22.9	0.6	4.8	
1	2008-12-02	1	7.4	25.1	0.0	4.8	
2	2008-12-03	1	12.9	25.7	0.0	4.8	
3	2008-12-04	1	9.2	28.0	0.0	4.8	
4	2008-12-05	1	17.5	32.3	1.0	4.8	
...	...	...	...	...	...	...	
145455	2017-06-21	49	2.8	23.4	0.0	4.8	
145456	2017-06-22	49	3.6	25.3	0.0	4.8	
145457	2017-06-23	49	5.4	26.9	0.0	4.8	
145458	2017-06-24	49	7.8	27.0	0.0	4.8	
145459	2017-06-25	49	14.9	22.6	0.0	4.8	

	Sunshine	WindGustDir	WindGustSpeed	WindDir9am	...	Pressure3pm	\
0	7.611178	1	44.0	1	...	1007.1	
1	7.611178	2	44.0	2	...	1007.8	
2	7.611178	3	46.0	1	...	1008.7	
3	7.611178	4	24.0	3	...	1012.8	
4	7.611178	1	41.0	4	...	1006.0	
...	...	...	...	...	...	...	
145455	7.611178	15	31.0	3	...	1020.3	
145456	7.611178	5	22.0	3	...	1019.1	
145457	7.611178	6	37.0	3	...	1016.8	
145458	7.611178	13	28.0	6	...	1016.5	
145459	7.611178	1	39.0	12	...	1017.9	

	Cloud9am	Cloud3pm	Temp9am	Temp3pm	RainToday	RainTomorrow	Year	\
0	8.000000	4.50993	16.9	21.8	0	0	2008	
1	4.447461	4.50993	17.2	24.3	0	0	2008	
2	4.447461	2.00000	21.0	23.2	0	0	2008	
3	4.447461	4.50993	18.1	26.5	0	0	2008	
4	7.000000	8.00000	17.8	29.7	0	0	2008	
...	...	...	...	...	...	...	...	
145455	4.447461	4.50993	10.1	22.4	0	0	2017	
145456	4.447461	4.50993	10.9	24.5	0	0	2017	
145457	4.447461	4.50993	12.5	26.1	0	0	2017	
145458	3.000000	2.00000	15.1	26.0	0	0	2017	
145459	8.000000	8.00000	15.0	20.9	0	0	2017	

	Month	Day
0	12	1
1	12	2
2	12	3
3	12	4
4	12	5
...	...	...
145455	6	21
145456	6	22
145457	6	23
145458	6	24
145459	6	25

[145460 rows x 26 columns]

In [183... df.dtypes

```
Out[183...] Date          datetime64[ns]
          Location      int64
          MinTemp       float64
          MaxTemp       float64
          Rainfall      float64
          Evaporation   float64
          Sunshine      float64
          WindGustDir    int64
          WindGustSpeed  float64
          WindDir9am     int64
          WindDir3pm     int64
          WindSpeed9am   float64
          WindSpeed3pm   float64
          Humidity9am    float64
          Humidity3pm    float64
          Pressure9am    float64
          Pressure3pm    float64
          Cloud9am       float64
          Cloud3pm       float64
          Temp9am        float64
          Temp3pm        float64
          RainToday      int64
          RainTomorrow   int64
          Year           int32
          Month          int32
          Day            int32
          dtype: object
```

```
In [184...] df.shape
```

```
Out[184...] (145460, 26)
```

```
In [185...] df.columns
```

```
Out[185...] Index(['Date', 'Location', 'MinTemp', 'MaxTemp', 'Rainfall', 'Evaporation',
                  'Sunshine', 'WindGustDir', 'WindGustSpeed', 'WindDir9am', 'WindDir3pm',
                  'WindSpeed9am', 'WindSpeed3pm', 'Humidity9am', 'Humidity3pm',
                  'Pressure9am', 'Pressure3pm', 'Cloud9am', 'Cloud3pm', 'Temp9am',
                  'Temp3pm', 'RainToday', 'RainTomorrow', 'Year', 'Month', 'Day'],
                  dtype='object')
```

## Outlier Detection

```
In [186...] # Assuming df is your DataFrame
# Replace the below line with your actual DataFrame loading or creation code
# df = pd.read_csv('your_data.csv')

# List of numeric columns for which outlier detection is suitable
numeric_columns = ['MinTemp', 'MaxTemp', 'Rainfall', 'Evaporation',
                  'Sunshine', 'WindGustSpeed', 'WindSpeed9am', 'WindSpeed3pm',
                  'Humidity9am', 'Humidity3pm', 'Pressure9am', 'Pressure3pm',
                  'Cloud9am', 'Cloud3pm', 'Temp9am', 'Temp3pm']

# Dictionary to store the percentage of outliers for each column
```

```

outlier_percentages = {}

# Loop through the numeric columns and calculate the percentage of outliers
for column in numeric_columns:
    Q1 = df[column].quantile(0.25)
    Q3 = df[column].quantile(0.75)
    IQR = Q3 - Q1
    lower_bound = Q1 - 1.5 * IQR
    upper_bound = Q3 + 1.5 * IQR

    # Counting outliers
    outliers = df[(df[column] < lower_bound) | (df[column] > upper_bound)]
    total_count = df[column].count()
    outlier_percentage = (len(outliers) / total_count) * 100
    outlier_percentages[column] = outlier_percentage

# Print the percentage of outliers for each column
for column, percentage in outlier_percentages.items():
    print(f'Percentage of outliers in {column}: {percentage:.2f}%')

```

```

Percentage of outliers in MinTemp: 0.06%
Percentage of outliers in MaxTemp: 0.37%
Percentage of outliers in Rainfall: 19.89%
Percentage of outliers in Evaporation: 25.76%
Percentage of outliers in Sunshine: 31.33%
Percentage of outliers in WindGustSpeed: 3.80%
Percentage of outliers in WindSpeed9am: 1.25%
Percentage of outliers in WindSpeed3pm: 1.73%
Percentage of outliers in Humidity9am: 0.98%
Percentage of outliers in Humidity3pm: 0.00%
Percentage of outliers in Pressure9am: 1.90%
Percentage of outliers in Pressure3pm: 1.74%
Percentage of outliers in Cloud9am: 0.00%
Percentage of outliers in Cloud3pm: 3.42%
Percentage of outliers in Temp9am: 0.21%
Percentage of outliers in Temp3pm: 0.68%

```



In [187...



```


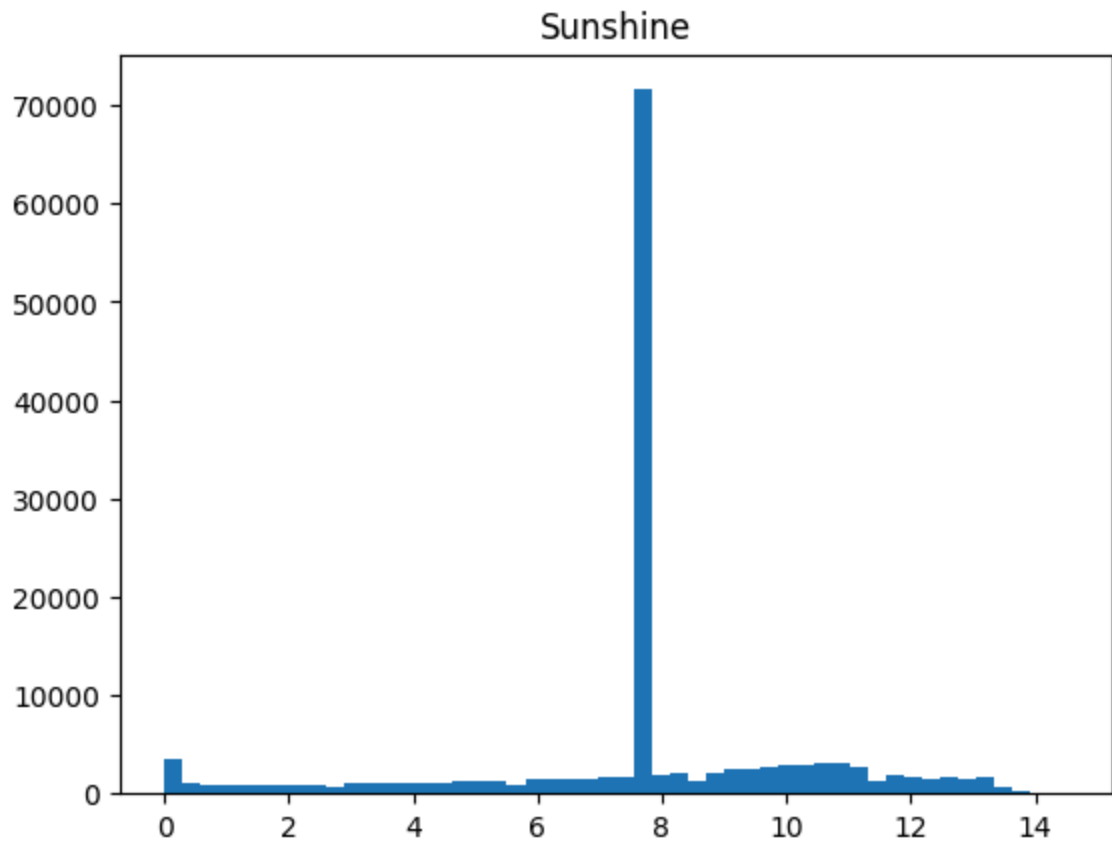
import matplotlib.pyplot as plt
numeric_columns = ['MinTemp', 'MaxTemp', 'Rainfall', 'Evaporation',
                   'Sunshine', 'WindGustSpeed', 'WindSpeed9am', 'WindSpeed3pm',
                   'Humidity9am', 'Humidity3pm', 'Pressure9am', 'Pressure3pm',
                   'Cloud9am', 'Cloud3pm', 'Temp9am', 'Temp3pm']



for col in numeric_columns:
    plt.hist(df[col], bins=50)
    plt.title(col)
    plt.show()



```



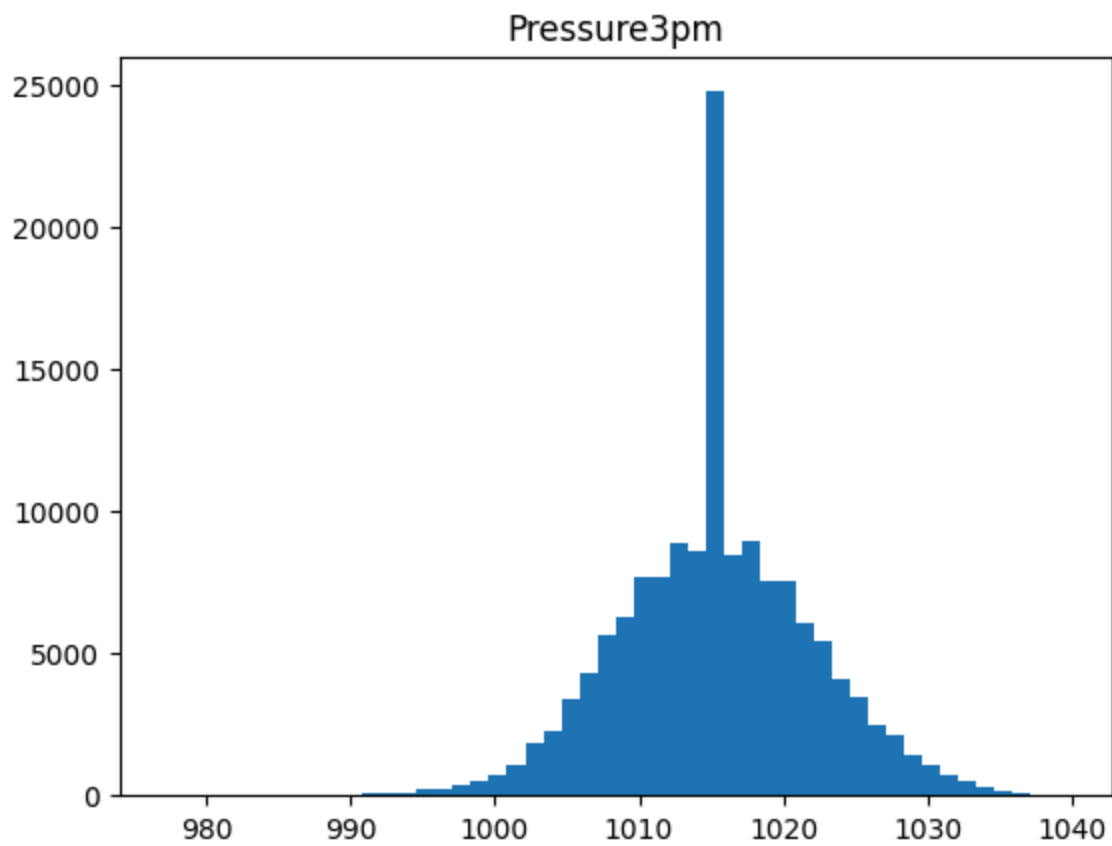







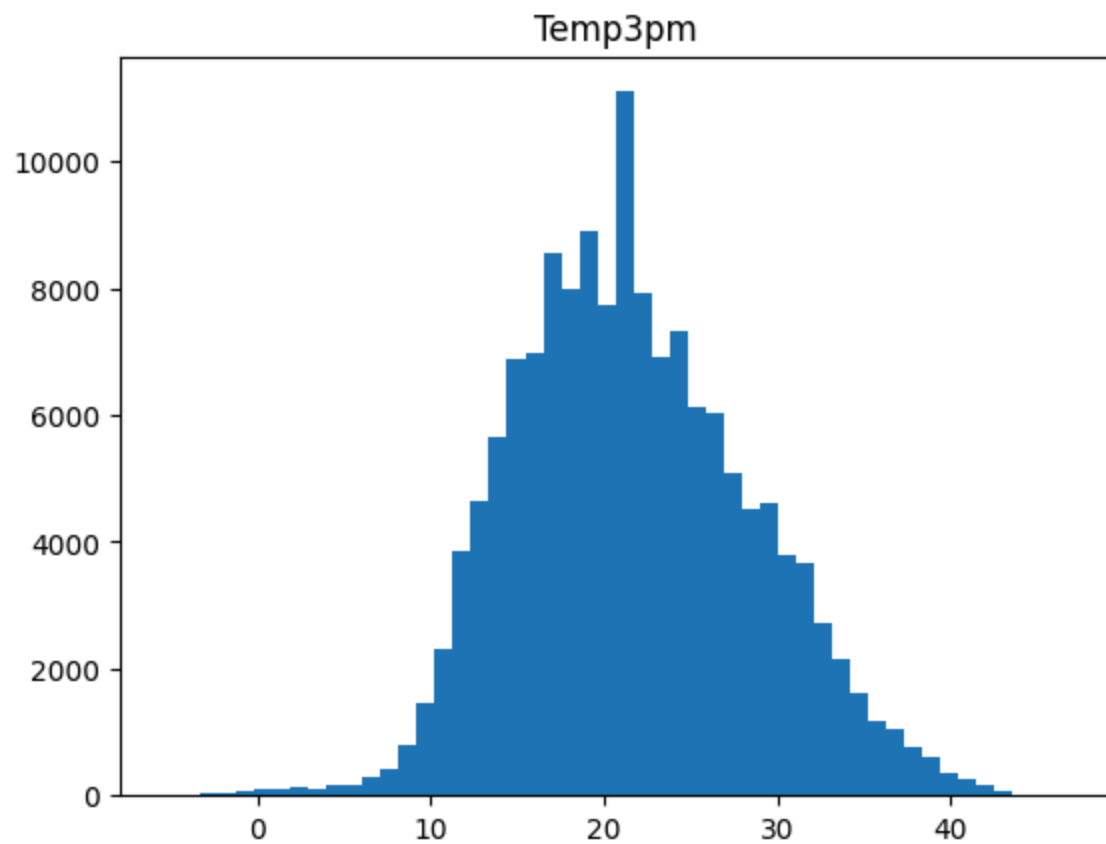













## Handling Outliers

In [188...

```
import pandas as pd

# Function to cap outliers
def cap_outliers(series, iqr_factor=1.5):
    Q1 = series.quantile(0.25)
    Q3 = series.quantile(0.75)
    IQR = Q3 - Q1
    lower_cap = Q1 - iqr_factor * IQR
    upper_cap = Q3 + iqr_factor * IQR
    return series.clip(lower=lower_cap, upper=upper_cap)

# Applying the function to each numeric column in the DataFrame
numeric_cols = df.select_dtypes(include=[np.number]).columns # Adjust as necessary
for col in numeric_cols:
    df[col] = cap_outliers(df[col])

# Optional: Check results for one of the columns or summary stats
print(df.describe()) # Provides summary statistics to check the max and min, ensur
```

	Date	Location	MinTemp \
count	145460	145460.000000	145460.000000
mean	2013-04-04 21:08:51.907053568	24.560126	12.192336
min	2007-11-01 00:00:00	1.000000	-5.950000
25%	2011-01-11 00:00:00	12.000000	7.700000
50%	2013-06-02 00:00:00	24.000000	12.000000
75%	2015-06-14 00:00:00	37.000000	16.800000
max	2017-06-25 00:00:00	49.000000	30.450000
std	NaN	13.941805	6.364499

	MaxTemp	Rainfall	Evaporation	Sunshine \
count	145460.000000	145460.000000	145460.000000	145460.000000
mean	23.219758	0.381674	4.750932	7.922535
min	2.700000	0.000000	2.200000	5.977944
25%	18.000000	0.000000	4.000000	7.611178
50%	22.600000	0.000000	4.800000	7.611178
75%	28.200000	0.600000	5.200000	8.700000
max	43.500000	1.500000	7.000000	10.333234
std	7.067804	0.608638	1.454089	1.386787

	WindGustDir	WindGustSpeed	WindDir9am ...	Pressure3pm \
count	145460.000000	145460.000000	145460.000000	145460.000000
mean	8.013028	39.64328	8.463151	1015.268537
min	1.000000	8.50000	1.000000	998.650000
25%	3.000000	31.00000	5.000000	1011.100000
50%	8.000000	39.00000	9.000000	1015.200000
75%	12.000000	46.00000	12.000000	1019.400000
max	16.000000	68.50000	16.000000	1031.850000
std	4.905515	12.17591	4.399079	6.528909

	Cloud9am	Cloud3pm	Temp9am	Temp3pm	RainToday \
count	145460.000000	145460.000000	145460.000000	145460.000000	145460.0
mean	4.447461	4.544125	16.988207	21.685669	0.0
min	0.000000	1.000000	-1.500000	2.450000	0.0
25%	3.000000	4.000000	12.300000	16.700000	0.0
50%	4.447461	4.509930	16.700000	21.400000	0.0
75%	6.000000	6.000000	21.500000	26.200000	0.0
max	9.000000	9.000000	35.300000	40.450000	0.0
std	2.265604	2.026092	6.440883	6.812734	0.0

	RainTomorrow	Year	Month	Day
count	145460.0	145460.000000	145460.000000	145460.000000
mean	0.0	2012.769751	6.399615	15.712258
min	0.0	2007.000000	1.000000	1.000000
25%	0.0	2011.000000	3.000000	8.000000
50%	0.0	2013.000000	6.000000	16.000000
75%	0.0	2015.000000	9.000000	23.000000
max	0.0	2017.000000	12.000000	31.000000
std	0.0	2.537684	3.427262	8.794789

[8 rows x 26 columns]

## Feature Engineering

```
In [189... df['TempChange9amTo3pm'] = df['Temp3pm'] - df['Temp9am']
df['HumidityChange9amTo3pm'] = df['Humidity3pm'] - df['Humidity9am']
df['PressureChange9amTo3pm'] = df['Pressure3pm'] - df['Pressure9am']
print(df[['TempChange9amTo3pm', 'HumidityChange9amTo3pm', 'PressureChange9amTo3pm']])
```

	TempChange9amTo3pm	HumidityChange9amTo3pm	PressureChange9amTo3pm
0	4.9	-49.0	-0.6
1	7.1	-19.0	-2.8
2	2.2	-8.0	1.1
3	8.4	-29.0	-4.8
4	11.9	-49.0	-4.8

## Model Development

```
In [190... import sklearn
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import accuracy_score
```

```
In [191... numeric_df = df.select_dtypes(include=['number'])


# Calculate the correlation matrix
corr = numeric_df.corr()

# Set up the matplotlib figure
plt.figure(figsize=(22, 10)) # Adjust the size as needed

# Generate a heatmap
sns.heatmap(corr, annot=True, fmt=".1f", cmap='coolwarm',
            cbar=True, square=True, linewidths=.5)

# Add a title to the heatmap
plt.title('Correlation Matrix of Numerical Features')

# Show the plot
plt.show()
```



In [192...

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.impute import SimpleImputer
from sklearn.ensemble import HistGradientBoostingClassifier
from sklearn.metrics import accuracy_score

# 2. Drop rows with missing values in RainTomorrow (target variable)
df.dropna(subset=["RainTomorrow"], inplace=True)

# 3. Separate features and target
X = df.drop(["RainTomorrow", "Date"], axis=1) # Drop "Date" column
y = df["RainTomorrow"]

# 4. Encode categorical features
X = pd.get_dummies(X)

# 5. Split data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# 6. Impute missing values using SimpleImputer
imputer = SimpleImputer(strategy="mean") # Replace missing values with the mean
X_train_imputed = imputer.fit_transform(X_train)
X_test_imputed = imputer.transform(X_test)
```

```
# 7. Train a model that supports missing values (e.g., HistGradientBoostingClassifier)
model = HistGradientBoostingClassifier()
model.fit(X_train_imputed, y_train)

# 8. Make predictions and evaluate
y_pred = model.predict(X_test_imputed)
accuracy = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy*100)
```

Accuracy: 100.0

## Parameter Tuning

In [193...

```
import pandas as pd
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import accuracy_score
param_grid = {
    'max_depth': [3, 5, 7, 10],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 2, 4]
}
model = DecisionTreeClassifier()
grid_search = GridSearchCV(estimator=model, param_grid=param_grid, cv=5, scoring='a
grid_search.fit(X_train, y_train)
print("Best Parameters:", grid_search.best_params_)
print("Best Score:", grid_search.best_score_)
best_model = grid_search.best_estimator_
y_pred = best_model.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy)
```

Best Parameters: {'max\_depth': 3, 'min\_samples\_leaf': 1, 'min\_samples\_split': 2}

Best Score: 1.0

Accuracy: 1.0

## Model Evaluation

### Performance Metrics

In [194...

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score

# ... (Load data, preprocess, train model as before) ...

# Train the model
model = DecisionTreeClassifier()
model.fit(X_train, y_train)

# Predict probabilities for ROC-AUC calculation
y_pred_proba = model.predict_proba(X_test_imputed)[: , 1]
```



```
# Calculate metrics
accuracy = accuracy_score(y_test, y_pred)
precision = precision_score(y_test, y_pred, pos_label='Yes')
recall = recall_score(y_test, y_pred, pos_label='Yes')
f1 = f1_score(y_test, y_pred, pos_label='Yes')
roc_auc = roc_auc_score(y_test, y_pred_proba)

print("Accuracy:", accuracy)
print("Precision:", precision)
print("Recall:", recall)
print("F1-score:", f1)
print("ROC-AUC:", roc_auc)
```

/usr/local/lib/python3.10/dist-packages/sklearn/base.py:439: UserWarning: X does not have valid feature names, but DecisionTreeClassifier was fitted with feature names  
warnings.warn(

```
-----
IndexError                                Traceback (most recent call last)
<ipython-input-194-d97cb17528a3> in <cell line: 13>()
    11
    12 # Predict probabilities for ROC-AUC calculation
----> 13 y_pred_proba = model.predict_proba(X_test_imputed)[: , 1]
    14
    15 # Calculate metrics

IndexError: index 1 is out of bounds for axis 1 with size 1
```

## Validation Strategy

```
In [ ]: import pandas as pd
from sklearn.model_selection import train_test_split, KFold, cross_val_score
from sklearn.impute import SimpleImputer
from sklearn.ensemble import HistGradientBoostingClassifier

# ... (Load data, preprocess, define model as before) ...

# 1. Create KFold object
kfold = KFold(n_splits=10, shuffle=True, random_state=42) # 10 folds

# 2. Perform cross-validation
cv_scores = cross_val_score(model, X, y, cv=kfold, scoring='accuracy')

# 3. Print results
print("Cross-Validation Scores:", cv_scores)
print("Average Accuracy:", cv_scores.mean())
```

## Insights and Recommendation

```
In [ ]: import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.impute import SimpleImputer
from sklearn.ensemble import HistGradientBoostingClassifier
```

```
# ... (Load data, preprocess, train model as before) ...

# Get feature importances
importances = model.feature_importances_

# Get feature names
feature_names = X.columns

# Create a DataFrame with feature names and importances
feature_importance_df = pd.DataFrame({'Feature': feature_names, 'Importance': importances})

# Sort by importance in descending order
feature_importance_df = feature_importance_df.sort_values(by='Importance', ascending=False)

# Print the DataFrame
print(feature_importance_df)
```

In [ ]: